

# Playbook for Verizon Migration of Cloud Pak for Data from NFS to ODF

---

## Runbook Version

Runbook Version: **1.4.0**

## Table of Contents

---

- [Cloud Pak for Data operators backup and restore](#)
  - [MTC installation and setup on source and target clusters](#)
  - [MTC migration plan creation](#)
  - [Shutdown Cloud Pak for Data instance on source cluster](#)
  - [MTC migration stage to target cluster](#)
  - [MTC migration cutover to target cluster](#)
  - [Patch and update new storage class references on target cluster](#)
  - [Start Cloud Pak for Data instance on target cluster](#)
- 

This migration playbook requires the `cpd-migrate.sh` and `cpd-migrate.input.json` (also referenced as `input.json`) which will also be provided to Verizon alongside this playbook.

Additionally, this playbook is for clusters on a restricted network. Below are the links for offline installations of CPDBR, MTC and OADP.

CPDBR: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=utility-moving-images-backup-restore-utilities>

MTC: [https://access.redhat.com/documentation/en-us/openshift\\_container\\_platform/4.10/html/migration\\_toolkit\\_for\\_containers/installing-mtc-restricted](https://access.redhat.com/documentation/en-us/openshift_container_platform/4.10/html/migration_toolkit_for_containers/installing-mtc-restricted)

MTC rsync fix image: [quay.io/konveyor/mig-controller@sha256:68effcbddddd6a96c3185daea84f8741e6881fb117ca2e70d76957d53ae779c79](https://quay.io/repository/konveyor/mig-controller@sha256:68effcbddddd6a96c3185daea84f8741e6881fb117ca2e70d76957d53ae779c79)

OADP: <https://docs.openshift.com/container-platform/4.10/operators/admin/olm-restricted-networks.html#olm-restricted-networks>

## Cloud pak for data operators backup and restore

---

### OADP installation and setup

Reference: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=iobru-installing-cloud-pak-data-oadp-backup-restore-utility-components>

**This operation needs to be done on both the SOURCE cluster and the TARGET cluster.**

Follow steps in the official documentation above which include:

- **IMPORTANT:** If OADP operator is already installed in oadp-operator namespace, reinstall it to refresh the oadp crds, due to conflict between CPD and MTC versions of OADP that are used.
- Create the oadp-operator project
- Annotate the oadp-operator project so that Restic pods can be scheduled on all nodes
- Install OADP operator
- Create Secret and environment variables for S3 backups
- Create OADP CR

### Example of CPD OADP operator installation (v1.1.4):

The screenshot shows the Red Hat OpenShift Container Platform OperatorHub interface. The left sidebar contains navigation links: Administrator, Home, Operators (selected), Installed Operators, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area is titled "Install Operator" and includes the following sections:

- Update channel:** Radio buttons for stable, stable-1.0, and stable-1.1 (selected).
- Installation mode:** Radio buttons for "All namespaces on the cluster (default)" (disabled) and "A specific namespace on the cluster" (selected). A note states: "This mode is not supported by this Operator" and "Operator will be available in a single Namespace only."
- Installed Namespace:** Radio buttons for "Operator recommended Namespace: openshift-adp" and "Select a Namespace". A dropdown menu shows "oadp-operator" (selected).
- Update approval:** Radio buttons for Automatic and Manual (selected).

A warning message states: "Not installing the Operator into the recommended namespace can cause unexpected behavior." A blue information box notes: "Manual approval applies to all operators in a namespace. Installing an operator with manual approval causes all operators installed in namespace oadp-operator to function as manual approval strategy. To allow automatic approval, all operators installed in the namespace must use automatic approval strategy."

On the right, the "OADP Operator" is listed as provided by Red Hat. Under "Provided APIs", the following are listed:

- Backup:** Backup is a Velero resource that represents the capture of Kubernetes cluster state at a point in time (API objects and associated volume state).
- BSL BackupStorageLocation:** BackupStorageLocation represents an object storage location (such as Amazon S3 Bucket) where Velero stores backup objects.
- DBR DeleteBackupRequest:** DeleteBackupRequest is a request to delete one or more backups.
- DR DownloadRequest:** DownloadRequest is a request to download an artifact from object storage.

```
$ oc get csv -n oadp-operator
NAME                                DISPLAY          VERSION  REPLACES
PHASE
oadp-operator.v1.1.4  OADP Operator   1.1.4    oadp-operator.v1.1.3
Succeeded
```

### Example of creation of s3 object storage credentials secret:

In this example, the s3 object storage credential secret is called `cloud-credentials`:

```
$ cat credentials-velero-2
[default]
aws_access_key_id = 131f141055654845adf3a7918178xxxx
aws_secret_access_key = xxxx
$ oc create secret generic cloud-credentials \
```

```
--namespace oadp-operator \
--from-file cloud=./credentials-velero-2
secret/cloud-credentials created
$
```

**Example of creation of OADP CR. In this example it is called `dataprotectionapp-2.yaml`.**

```
$ cat dataprotectionapp-2.yaml
apiVersion: oadp.openshift.io/v1alpha1
kind: DataProtectionApplication
metadata:
  name: cpst-dpa
spec:
  configuration:
    velero:
      customPlugins:
        - image: icr.io/cpopen/cpd/cpddb-velero-plugin:4.0.0-beta1-1-x86_64
          name: cpddb-velero-plugin
      defaultPlugins:
        - aws
        - openshift
        - csi
      podConfig:
        resourceAllocations:
          limits:
            cpu: "1"
            memory: 1Gi
          requests:
            cpu: 500m
            memory: 256Mi
    restic:
      enable: true
      timeout: 12h
      podConfig:
        resourceAllocations:
          limits:
            cpu: "1"
            memory: 8Gi
          requests:
            cpu: 500m
            memory: 256Mi
          tolerations:
            - key: icp4data
              operator: Exists
              effect: NoSchedule
  backupImages: false
  backupLocations:
    - velero:
        provider: aws
        default: true
        objectStorage:
          bucket: mtc-todd
```

```
    prefix: cpst-9c9f-backup
  config:
    region: us-south
    s3ForcePathStyle: "true"
    s3Url: https://s3.us-south.cloud-object-storage.appdomain.cloud
  credential:
    name: cloud-credentials
    key: cloud
$ oc apply -f dataprotectionapp-2.yaml
dataprotectionapplication.oadp.openshift.io/cpst-dpa created
$
```

## Create cpd operators configmap on source cluster

Reference: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=utility-downloading-backup-restore-scripts>

Download cpd-operators.sh script:

```
wget https://raw.githubusercontent.com/IBM/cpd-
cli/master/cpdops/files/cpd-operators.sh
```

Note: foundation-namespace and operators namespace are the same for cpd express installations (ibm-common-services by default)

```
cpd-operators.sh backup --foundation-namespace ibm-common-services --
operators-namespace ibm-common-services
```

**Verify the configmap is present:**

```
oc get configmap cpd-operators -n ibm-common-services
```

## Backup cpd operators on source cluster

Reference: [https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=obr-scenario-creating-offline-backup-cloud-pak-data-instance-restoring-it-different-cluster#oadp-new-cluster\\_\\_backup\\_foundational\\_operators](https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=obr-scenario-creating-offline-backup-cloud-pak-data-instance-restoring-it-different-cluster#oadp-new-cluster__backup_foundational_operators)

**Configure cpd-cli oadp to point to the project where the Velero instance is installed:**

```
cpd-cli oadp client config set namespace=oadp-operator
```

**Trigger a cpd oadp offline backup for cpd-operators namespace:**

```
cpd-cli oadp backup create <cpd-ops-backup-name> --include-namespaces ibm-  
common-services --include-  
resources='namespaces,operatorgroups,configmaps,scheduling,crd' --skip-  
hooks --log-level=debug --verbose
```

**List backups:**

```
cpd-cli oadp backup ls
```

**Verify the backup:**

```
cpd-cli oadp backup status <cpd-ops-backup-name> --details
```

## Restore cpd operators on target cluster

**Check catalog sources on the target cluster and clean up any CPD related catalog sources**

Before running the CPD operator subscriptions restore, make sure the catalog sources are cleaned of CPD related ones on the target cluster.

Example output, showing no CPD related catalog sources:

```
$ oc get catsrc -n openshift-marketplace  
NAME                                DISPLAY                                TYPE    PUBLISHER    AGE  
certified-operators                Certified Operators                    grpc    Red Hat      21d  
community-operators                Community Operators                    grpc    Red Hat      21d  
ibm-operator-catalog               IBM Operator Catalog                  grpc    IBM          21d  
redhat-marketplace                 Red Hat Marketplace                    grpc    Red Hat      21d  
redhat-operators                   Red Hat Operators                      grpc    Red Hat      21d  
$
```

Verify ibm-common-services namespace does not exist:

```
oc get ns |grep common
```

**Configure cpd-cli oadp to point to the project where the Velero instance is installed:**

```
cpd-cli oadp client config set namespace=oadp-operator
```

**Restore cpd operators crds:**

```
cpd-cli oadp restore create <cpd-operators-restore-crds> \
--from-backup=<cpd-ops-backup-name> \
--include-resources='crd' \
--include-cluster-resources=true \
--skip-hooks \
--log-level=debug \
--verbose
```

#### Verify crd restore operation:

```
cpd-cli oadp restore status <cpd-operators-restore-crds> --details
```

#### Restore additional resources, such as projects and operator groups, etc:

Restore command:

```
cpd-cli oadp restore create <cpd-operators-restore-remaining> \
--from-backup=<cpd-ops-backup-name> \
--include-resources='namespaces,operatorgroups,scheduling,crd' \
--include-cluster-resources=true \
--skip-hooks \
--log-level=debug \
--verbose
```

#### Verify ibm-common-services namespace and resources were created:

```
$ oc get ns |grep ibm-common-services
ibm-common-services                                Active    67s
$ oc get operatorgroups -n ibm-common-services
NAME                AGE
operatorgroup       83s
```

#### Verify restore operation for remaining resources:

```
cpd-cli oadp restore status <cpd-operators-restore-remaining> --details
```

#### Restore cpd operators configmaps:

Configmaps before restore:

```
$ oc get cm -n ibm-common-services
NAME                DATA    AGE
```

kube-root-ca.crt	1	2m11s
openshift-service-ca.crt	1	2m11s

Restore command:

```
cpd-cli oadp restore create <cpd-operators-restore-cm> \
--from-backup=<cpd-ops-backup-name> \
--include-resources='configmaps' \
--selector 'app=cpd-operators-backup' \
--skip-hooks \
--log-level=debug \
--verbose
```

Configmaps after restore:

```
$ oc get cm -n ibm-common-services
NAME                DATA  AGE
cpd-operators        18     24s
kube-root-ca.crt     1      3m30s
openshift-service-ca 1      3m30s
```

**Verify restore operation for configmaps:**

```
cpd-cli oadp restore status <cpd-operators-restore-cm>
```

**Restore operator subscriptions** Reference: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=utility-downloading-backup-restore-scripts>

**Download cpd-operators.sh script:**

```
wget https://raw.githubusercontent.com/IBM/cpd-
cli/master/cpdops/files/cpd-operators.sh
```

**Run command:**

```
cpd-operators.sh restore --foundation-namespace ibm-common-services --
operators-namespace ibm-common-services
```

**Verify restore operator subscriptions:**

```

$ oc get sub -n ibm-common-services
NAME
PACKAGE                                SOURCE
CHANNEL
cpd-operator
cpd-platform-operator                  cpd-platform
v3.8
fdb-kubernetes-operator-v2.4-ibm-fdb-operator-catalog-openshift-
marketplace    fdb-kubernetes-operator          ibm-fdb-operator-
catalog                                v2.4
ibm-cert-manager-operator
ibm-cert-manager-operator              opencloud-operators
v3.23
ibm-common-service-operator
ibm-common-service-operator            opencloud-operators
v3.23
ibm-cpd-ae-operator
analyticsengine-operator                ibm-cpd-ae-operator-catalog
v3.5
ibm-cpd-ccs-operator
ibm-cpd-ccs                            ibm-cpd-ccs-operator-catalog
v6.5
ibm-cpd-datarefinery-operator
ibm-cpd-datarefinery                    ibm-cpd-datarefinery-operator-catalog
v6.5
ibm-cpd-datastage-operator
ibm-cpd-datastage-operator              ibm-cpd-datastage-operator-catalog
v3.5
ibm-cpd-iis-operator
ibm-cpd-iis                            ibm-cpd-iis-operator-catalog
v3.5
ibm-cpd-wkc-operator-catalog-subscription
ibm-cpd-wkc                            ibm-cpd-wkc-operator-catalog
v3.5
ibm-cpd-wml-operator
ibm-cpd-wml-operator                    ibm-cpd-wml-operator-catalog
v3.5
ibm-cpd-ws-operator
ibm-cpd-wsl                            ibm-cpd-ws-operator-catalog
v6.5
ibm-cpd-ws-runtimes-operator
ibm-cpd-ws-runtimes                    ibm-cpd-ws-runtimes-operator-catalog
v6.5
ibm-db2aaservice-cp4d-operator
ibm-db2aaservice-cp4d-operator          ibm-db2aaservice-cp4d-operator-
catalog                                v3.2
ibm-db2oltp-cp4d-operator-catalog-subscription
ibm-db2oltp-cp4d-operator                ibm-db2oltp-cp4d-operator-catalog
v3.2
ibm-db2u-operator
db2u-operator                          ibm-db2uoperator-catalog
v3.2
ibm-db2wh-cp4d-operator-catalog-subscription

```



ibm-db2wh-cp4d-operator v3.2	ibm-db2wh-cp4d-operator-catalog
ibm-dmc-operator-subscription ibm-dmc-operator v2.2	ibm-dmc-operator-catalog
ibm-fdb-operator ibm-opencontent-foundationdb v2.4	ibm-fdb-operator-catalog
ibm-namespace-scope-operator ibm-namespace-scope-operator v3.23	opencloud-operators
ibm-watson-openscale-operator-subscription ibm-cpd-wos v3.5	ibm-openscale-operator-catalog
ibm-zen-operator ibm-zen-operator v3.23	opencloud-operators
manta-adl-operator manta-adl-operator v1.10	manta-adl-operator-catalog
operand-deployment-lifecycle-manager-app ibm-odlm v3.23	opencloud-operators
redis-operator ibm-cloud-databases-redis-operator catalog v1.6	ibm-cloud-databases-redis-operator-catalog

```
$ oc get csv -n ibm-common-services
```

NAME	DISPLAY
VERSION REPLACES	PHASE
cpd-platform-operator.v3.8.0	Cloud Pak for Data Platform
Operator	3.8.0 cpd-platform-
operator.v3.7.0 Succeeded	
db2u-operator.v3.2.0	IBM Db2
3.2.0 Succeeded	
fdb-kubernetes-operator.v2.4.6	FoundationDB Kubernetes
2.4.6 Succeeded	
ibm-cert-manager-operator.v3.25.2	IBM Cert Manager
3.25.2 Succeeded	
ibm-cloud-databases-redis.v1.6.6	IBM Operator for Redis
1.6.6 ibm-cloud-databases-redis.v1.6.5 Succeeded	
ibm-common-service-operator.v3.23.2	IBM Cloud Pak foundational
services 3.23.2 Succeeded	
ibm-cpd-ae.v3.5.0	IBM Analytics Engine
Powered by Apache Spark Service 3.5.0 Succeeded	
ibm-cpd-ccs.v6.5.0	Common Core Services
6.5.0 Succeeded	
ibm-cpd-datarefinery.v6.5.0	IBM Data Refinery
6.5.0 Succeeded	

ibm-cpd-datastage-operator.v3.5.0	IBM DataStage
3.5.0	Succeeded
ibm-cpd-iis.v1.6.5	IIS Services
1.6.5	Succeeded
ibm-cpd-wkc.v1.6.5	WKC Services
1.6.5	Succeeded
ibm-cpd-wml-operator.v3.5.0	IBM WML Services
3.5.0	Succeeded
ibm-cpd-wos.v3.5.0	IBM Watson OpenScale
3.5.0	Succeeded
ibm-cpd-ws-runtimes.v6.5.0	Watson Studio Notebook
Runtimes	6.5.0
Succeeded	
ibm-cpd-wsl.v6.5.0	Watson Studio
6.5.0	Succeeded
ibm-databases-dmc.v2.2.0	IBM Db2 Data Management
Console	2.2.0
Succeeded	
ibm-db2aaservice-cp4d-operator.v3.2.0	IBM Db2 as a Service
Operator Extension for IBM Cloud Pak for Data	3.2.0
Succeeded	
ibm-db2oltp-cp4d-operator.v3.2.0	IBM® Db2 Operator Extension
for IBM Cloud Pak for Data	3.2.0
Succeeded	
ibm-db2wh-cp4d-operator.v3.2.0	IBM® Db2 Warehouse Operator
Extension for IBM® Cloud Pak for Data	3.2.0
Succeeded	
ibm-namespace-scope-operator.v1.17.2	IBM NamespaceScope Operator
1.17.2	Succeeded
ibm-opencontent-foundationdb.v2.4.6	IBM Opencontent
FoundationDB	2.4.6
Succeeded	
ibm-zen-operator.v1.8.3	IBM Zen Service
1.8.3	Succeeded
manta-adl-operator.v1.10.0	MANTA Automated Data
Lineage	1.10.0
Succeeded	
operand-deployment-lifecycle-manager.v1.21.2	Operand Deployment
Lifecycle Manager	1.21.2
Succeeded	

Note: The exact subscription and csv installed can vary based on the services installed on source cluster.

## Collect an oadp resources dump of the cpd namespace on source cluster

To collect CPD resources from source cluster, perform the following:

```
cpd-cli oadp backup create <src-cpd-instance-resources> \
--include-namespaces <cpd-namespace> \
--exclude-resources='event,event.events.k8s.io,image-tags.openshift.io' \
--include-cluster-resources=true \
```

```
--snapshot-volumes=false \  
--skip-hooks=true \  
--log-level=debug \  
--verbose
```

Download the resource list:

```
cpd-cli oadp backup download <src-cpd-instance-resources> -o <src-cpd-  
instance-resources>-data.tar.gz
```

Provide the data.tar.gz tarball to IBM. This will be used to help identify all of the resources with storage class references and will be used to build the input.json file, which will be used in a later step to run the cpd-migrate.sh tool.

### Example for how to use the cpd resource tarball to identify resources with storage class references:

Extract tarball

```
$ tar -xzf src-cpd-instance-resources-data.tar.gz  
$ pwd  
/home/admin/cp4d/cpd-cli-linux-EE-12.0.5-61/src-cpd-instance-resources-  
data  
$ ls -l  
total 12  
drwxrwxr-x.  2 admin admin   21 May 15 19:44 metadata  
drwxrwxr-x. 129 admin admin 8192 May 15 19:44 resources
```

From the extracted oadp tar bundle, we can perform a series of grep commands to narrow in on the subset of k8s resources that are affected with storage class references.

Use helper commands to find potential resources to change.

- If looking at the cpd-cli oadp backup tar bundle, example grep commands to help find potential resources for patching.
- The best grep command will be to grep on the actual storageclass name(s). This will directly find where the storageclass is referenced.

```
grep -Rl -E "<cpd-storage-class>" ./* | \  
grep -vE  
"persistentvolumes|persistentvolumeclaims|preferredversion|storageclasses.  
storage.k8s.io"
```

- The next grep command is to identify which of those json files from the previous grep might have both RWO and RWX references in them.

```
grep -Rl -E "<cpd-storage-class>" ./* | \
grep -vE
"persistentvolumes|persistentvolumeclaims|preferredversion|storageclasses.
storage.k8s.io" | \
xargs -I{} sh -c "echo ----- {} -----;cat {}" | \
jq . | grep -E 'ReadWriteOnce|ReadWriteMany|<cpd-storage-class>' "
```

## MTC installation and setup on source and target clusters

### Summary of steps in this section:

- **IMPORTANT:** If MTC is already installed in openshift-migration namespace, reinstall just the OADP operator (version v1.0.x) in the openshift-migration namespace to refresh the oadp crds, due to conflict between CPD and MTC versions of OADP that are used.
- Install Migration Toolkit for Containers Operator (MTC) (currently v1.7.9) on both source and target clusters
- From source cluster MTC, add target cluster "Add cluster"
- Add replication repository (s3 object store, accessible from both source and target clusters)

### Install MTC and create MigrationController on source and target clusters

MTC and MigrationController need to be installed and created on both the source cluster and the target cluster.

Install MTC operator from Operator hub:

The screenshot displays the Red Hat OpenShift OperatorHub interface. On the left, a sidebar shows navigation options like Administrator, Home, Overview, Projects, Search, API Explorer, Events, Operators, Installed Operators, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area shows the 'OperatorHub' page with a search bar and a list of operators. The 'Migration Toolkit for Containers Operator' is highlighted, showing its details. The details panel on the right includes the following information:

- Latest version:** 1.7.9
- Capability level:** Basic Install (selected), Seamless Upgrades, Full Lifecycle, Deep Insights, Auto Pilot.
- Source:** Red Hat
- Provider:** Red Hat
- Infrastructure features:** Disconnected
- Valid Subscriptions:** OpenShift Container Platform, OpenShift Platform Plus
- Repository:** <https://github.com/konveyor/mig-operator>
- Container image:** registry.redhat.io/rhmtc/openshift-migration-rhel8-operator:sha256:141d6c2c9a348c2bd5ec08e63bb23d3e214143c7dbae213bacb5ae471f7b8e3b

The screenshot shows the Red Hat OpenShift Container Platform console. The left sidebar contains navigation links: Administrator, Home, Operators (selected), Installed Operators, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area is titled 'Install Operator' and shows the configuration for the 'Migration Toolkit for Containers Operator'. The 'Update channel' is set to 'release-v1.7'. The 'Installation mode' is set to 'A specific namespace on the cluster'. The 'Installed Namespace' is 'openshift-migration'. A warning message states: 'Namespace monitoring: Please note that installing non-Red Hat operators into OpenShift namespaces and enabling monitoring voids user support. Enabling cluster monitoring for non-Red Hat operators can lead to malicious metrics data overriding existing cluster metrics. For more information, see the cluster monitoring documentation.' The 'Update approval' is set to 'Manual'. On the right, a list of 'Provided APIs' is shown, including DVM, DVMP, DIM, DISM, and MA.

After operator is installed, select to create a MigrationController instance.

- Change the view from "Form view" to "YAML view"
- Add the following properties to YAML for known rsync issue under the **spec** section:

```
mig_controller_image_fqin: quay.io/konveyor/mig-
controller@sha256:68effcbddd6a96c3185daea84f8741e6881fb117ca2e70d76957d53
ae779c79
migration_rsync_privileged: true
```

Example YAML with updated **spec** section:

```
apiVersion: migration.openshift.io/v1alpha1
kind: MigrationController
metadata:
  creationTimestamp: '2023-05-25T07:55:19Z'
  generation: 2
  managedFields:
    - apiVersion: migration.openshift.io/v1alpha1
      fieldsType: FieldsV1
      fieldsV1:
        'f:spec':
```

```

    'f:migration_controller': {}
    'f:migration_log_reader': {}
    'f:cluster_name': {}
    'f:restic_timeout': {}
    'f:migration_rsync_privileged': {}
    'f:mig_pv_limit': {}
    'f:migration_velero': {}
    .: {}
    'f:mig_namespace_limit': {}
    'f:mig_controller_image_fqin': {}
    'f:azure_resource_group': {}
    'f:mig_pod_limit': {}
    'f:migration_ui': {}
    'f:olm_managed': {}
  manager: Mozilla
  operation: Update
  time: '2023-05-25T07:55:19Z'
- apiVersion: migration.openshift.io/v1alpha1
  fieldsType: FieldsV1
  fieldsV1:
    'f:status':
      .: {}
      'f:conditions': {}
  manager: ansible-operator
  operation: Update
  subresource: status
  time: '2023-05-25T07:55:19Z'
- apiVersion: migration.openshift.io/v1alpha1
  fieldsType: FieldsV1
  fieldsV1:
    'f:spec':
      'f:version': {}
  manager: OpenAPI-Generator
  operation: Update
  time: '2023-05-25T07:55:30Z'
name: migration-controller
namespace: openshift-migration
resourceVersion: '44253841'
uid: aa036929-1a16-4e7f-b967-5f0c3ad3bfb6
spec:
  mig_controller_image_fqin: >-
    quay.io/konveyor/mig-
controller@sha256:68effcbddddd6a96c3185daea84f8741e6881fb117ca2e70d76957d53
ae779c79
  mig_namespace_limit: '10'
  migration_ui: true
  mig_pod_limit: '100'
  migration_controller: true
  migration_log_reader: true
  olm_managed: true
  cluster_name: host
  restic_timeout: 1h
  migration_rsync_privileged: true
  migration_velero: true

```

```

mig_pv_limit: '100'
version: 1.7.9
azure_resource_group: ''
status:
conditions:
  - lastTransitionTime: '2023-05-25T07:56:16Z'
    message: ''
    reason: ''
    status: 'False'
    type: Failure
  - ansibleResult:
      changed: 0
      completion: '2023-05-31T21:17:47.506537'
      failures: 0
      ok: 51
      skipped: 22
    lastTransitionTime: '2023-05-25T07:55:19Z'
    message: Awaiting next reconciliation
    reason: Successful
    status: 'True'
    type: Running
  - lastTransitionTime: '2023-05-31T21:17:47Z'
    message: Last reconciliation succeeded
    reason: Successful
    status: 'True'
    type: Successful

```


Click "Create" and ensure MTC status is **Running**.

console-openshift-console.apps.ocp-120000pmky-9c9f.cloud.techzone.ibm.com/k8s/ns/openshift-migration/operators.coreos.com~v1alpha1~Clu...

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Project: openshift-migration

Installed Operators > Operator details

 **Migration Toolkit for Containers Operator**  
1.7.9 provided by Red Hat




Actions

Details YAML Subscription Events All instances DirectVolumeMigration DirectVolumeMigrationProgress DirectImageMigrat

### MigrationControllers

Create MigrationController

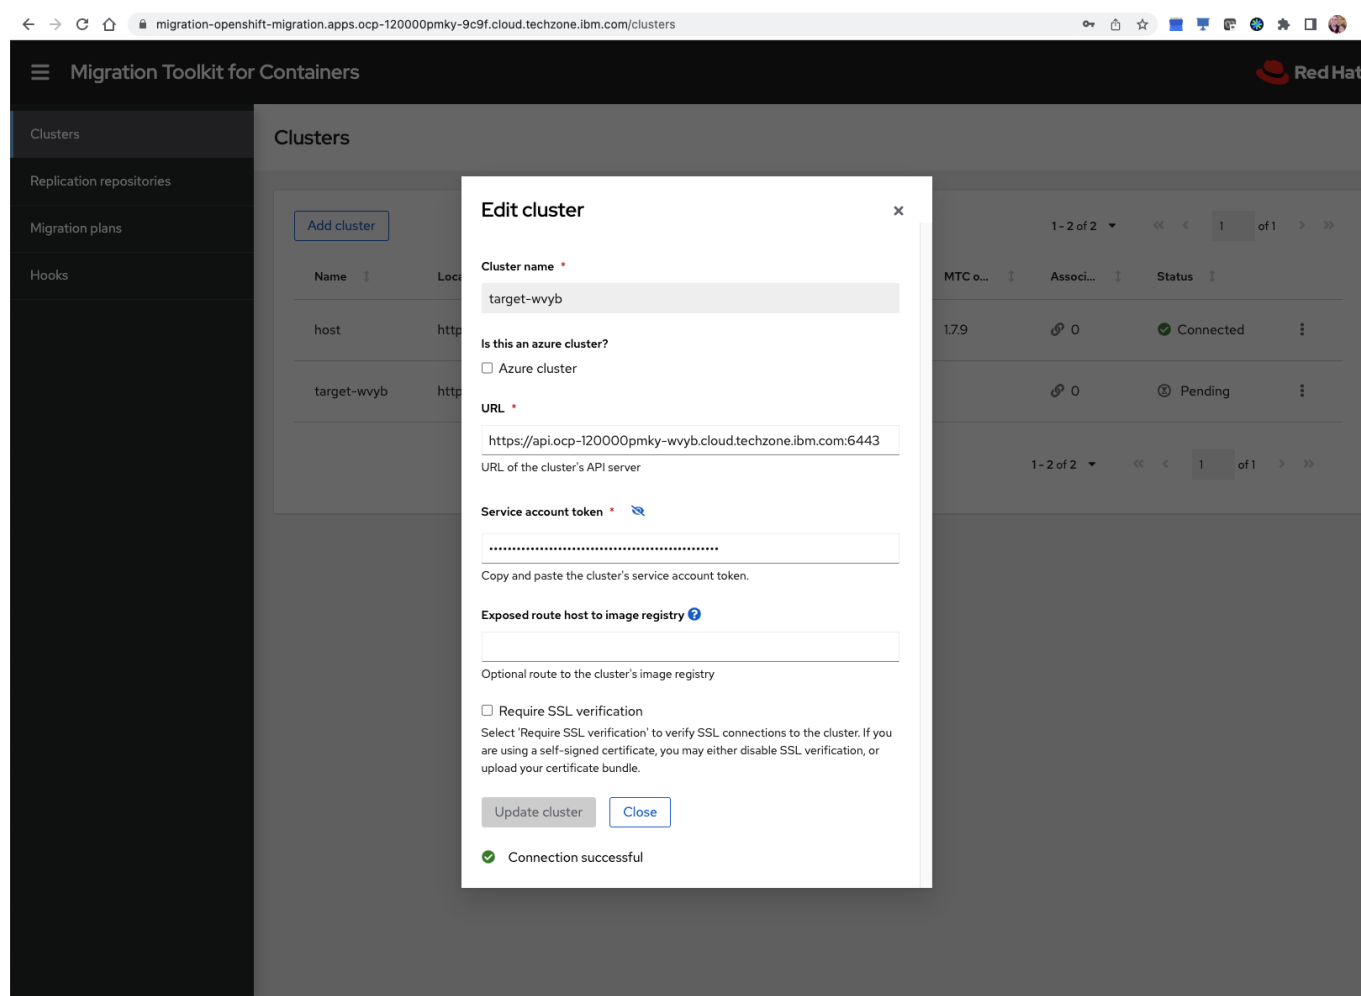
Name Search by name...

Name	Kind	Status	Labels	Last updated
 migration-controller	MigrationController	Condition:  Running	No labels	 Just now

Launch the migration controller console by clicking on the route created in "openshift-migration" namespace.

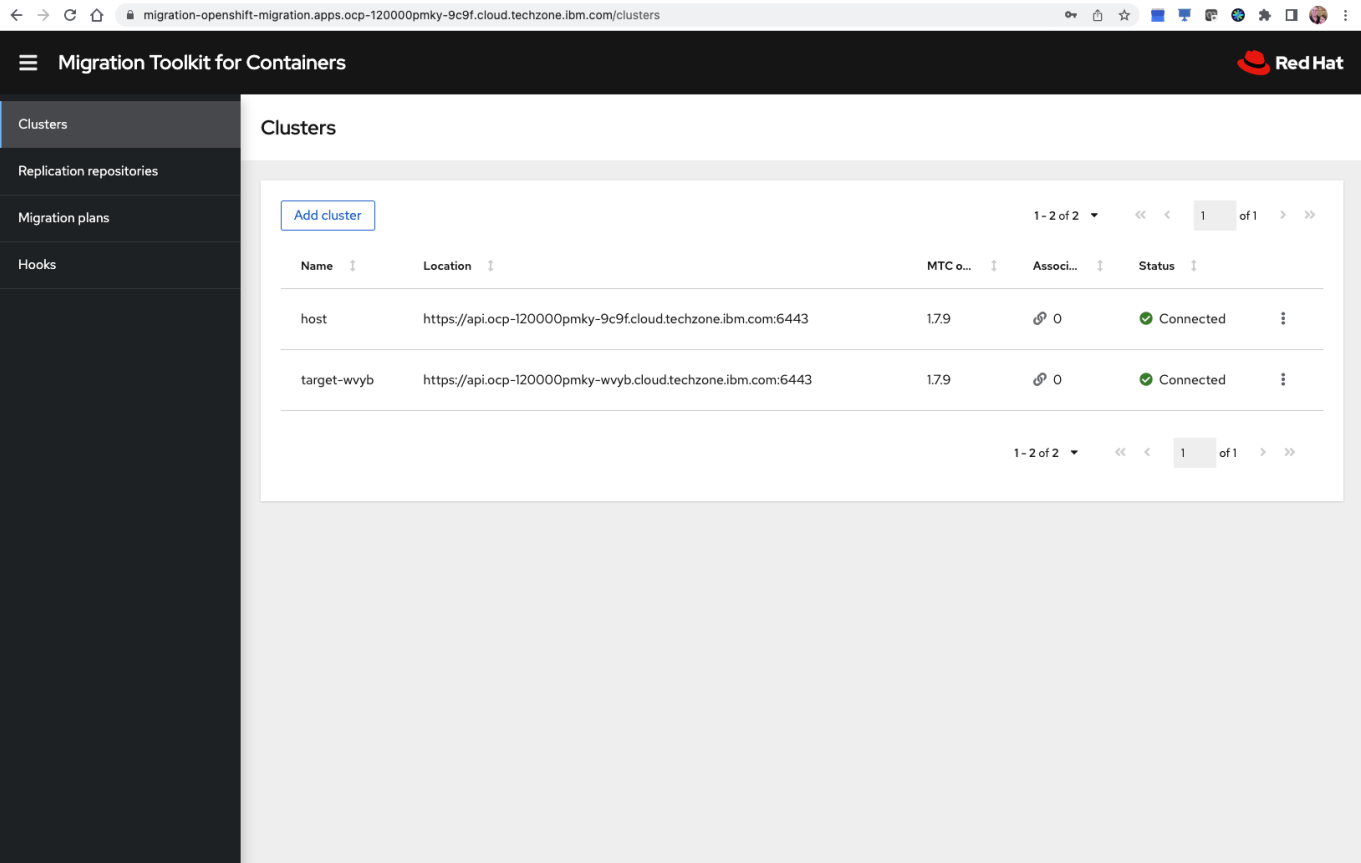
Add target cluster access details. For Service account token, you can execute the following commands in the target cluster. Please find the correct secret in your environment by executing `oc -n openshift-migration get secret|grep velero-token` and substitute it in the command below:

```
oc -n openshift-migration get secret velero-token-XXXXX -
ojsonpath='{.data.token}'|base64 -d
```



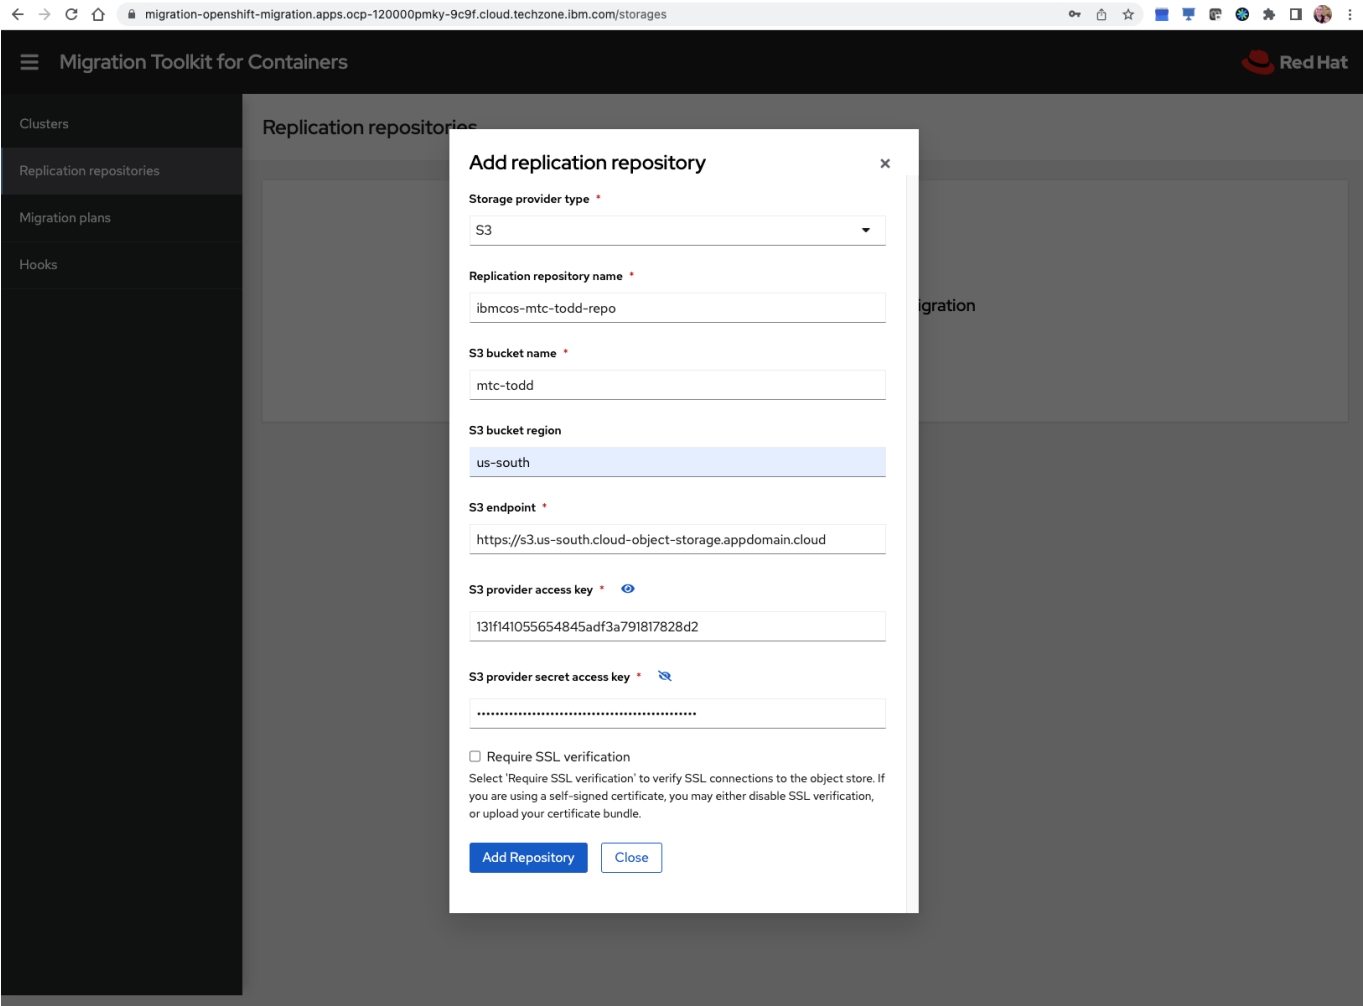


Next, verify source and target clusters are available and connected.

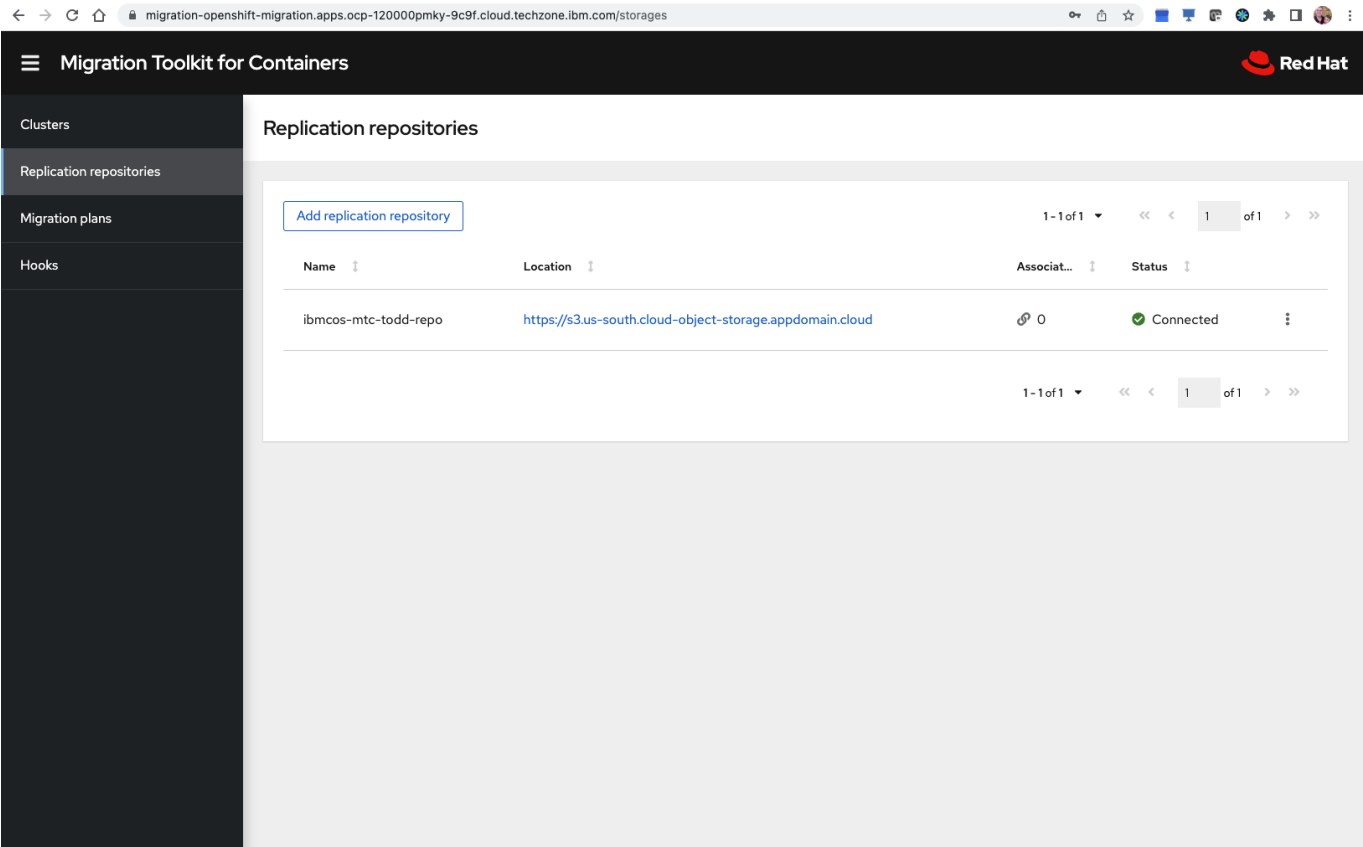


**Add replication repository (access to s3 object store):**

From source cluster MTC controller. Go to "Replication repositories" > "Add replication repository" Add the s3 object store access details and click "Add repository" and verify it is successfully connected.



Repo should now be accessible:



# MTC Migration plan creation

From source MTC migration console, select "Migration Plans" > "Add migration plan" Select "Full migration" migration type. Complete the remainder of the "General" section of the migration plan.

Create a migration plan

1 General

General

All fields are required.

Plan name \*

cp4d-465-wvyb-to-9c9f

Migration type \*

Full migration - migrate namespaces, persistent volumes (PVs) and Kubernetes resources from one cluster to another

Source cluster \*

host

Target cluster \*

target-wvyb

Repository \*

ibmcos-mtc-todd-repo

Next

Back

Cancel

**Namespaces:** Select your Cloud Pak for Data namespace.

Create a migration plan

1 General

2 Namespaces

3 Persistent volumes

4 Copy options

5 Migration options

6 Hooks

Namespaces

Select projects to be migrated.

Name

cpdinstance

Q

1 - 1 of 1

Name

cpdinstance

X

Clear all filters

<input type="checkbox"/>	Source name	Pods	PV claims	Services	Target name ?
<input checked="" type="checkbox"/>	cpdinstance	85	67	131	cpdinstance

1 selected

1 - 1 of 1

Next

Back

Cancel

**Persistent volumes:** For PVs, after persistent volumes are discovered, the MTC UI will have all PVs selected, with PV Migration Type as "Filesystem copy" by default. Leave the defaults and click "Next".

Create a migration plan

1 General

2 Namespaces

3 Persistent volumes

4 Copy options

5 Migration options

6 Hooks

Persistent volumes

Choose to move or copy persistent volumes associated with selected namespaces.

▼ PV name

Filter by PV name...

Q

1 - 10 of 67

<< < 1 of 7 > >>

<input checked="" type="checkbox"/>	PV name	Claim	Namespace	Storage class	Size	PV migration type	Details
<input checked="" type="checkbox"/>	pvc-087baf03-a71f-4b52-9ddd-d0f9aa594c3f	zookeeper-data-zookeeper-0	cpdinstance	cpd-storage	5Gi	Filesystem copy	<a href="#">View JSON</a>
<input checked="" type="checkbox"/>	pvc-0da3afcc-2a0e-4c67-9d98-0814bf277bd5	manta-configuration-service-claim	cpdinstance	cpd-storage	1Gi	Filesystem copy	<a href="#">View JSON</a>
<input checked="" type="checkbox"/>	pvc-0ef6bbda-laea-4b2a-b7ca-9db01d41a568	data-aiopyscale-ibm-aos-zookeeper-1	cpdinstance	cpd-storage	5Gi	Filesystem copy	<a href="#">View JSON</a>

Next

Back

Cancel

**Copy options:** Leave all of the default values and click "Next" without making any changes. In general, you would want to select the "file" storage class for RWX volumes and the "block" storage class for RWO volumes.

For ODF, it will be:  
RWX – file sc: ocs-storagecluster-cephfs  
RWO – block sc: ocs-storagecluster-ceph-rbd

However, MTC detects if the PVC is RWX vs RWO and attempts to select the correct storage class on the target cluster. Example screenshot:

Create a migration plan

1 General

2 Namespaces

3 Persistent volumes

4 Copy options

5 Migration options

6 Hooks

<input type="checkbox"/>	pvc-0da3afcc-2a0e-4c67-9d98-0814bf277bd5	manta-configuration-service-claim	cpd-storage	manta-configuration-service-claim	ocs-storagecluster-ceph-rbd.openshift-storage.rbd.csi.ceph.com	<input type="checkbox"/>
<input type="checkbox"/>	pvc-0ef6bbda-laea-4b2a-b7ca-9db01d41a568	data-aiopyscale-ibm-aos-zookeeper-1	cpd-storage	data-aiopyscale-ibm-aos-zookeeper-1	ocs-storagecluster-ceph-rbd.openshift-storage.rbd.csi.ceph.com	<input type="checkbox"/>
<input type="checkbox"/>	pvc-1003e7e9-f91d-4891-95c2-2eec177f0714	elasticsearch-master-elasticsearch-master-0	cpd-storage	elasticsearch-master-elasticsearch-master-0	ocs-storagecluster-ceph-rbd.openshift-storage.rbd.csi.ceph.com	<input type="checkbox"/>
<input type="checkbox"/>	pvc-108365f2-6a72-4e77-8cd1-93cf89f2f164	0072-iis-dedicatedservices-pvc	cpd-storage	0072-iis-dedicatedservices-pvc	ocs-storagecluster-cephfs.openshift-storage.cephfs.csi.ceph.com	<input type="checkbox"/>

Next

Back

Cancel

**Migration options:**

20 / 57

For Persistent volumes, the option "Use direct PV migration for filesystem copies" option has the following effect: If selected, then it will create an rsync server on the target cluster and perform rsync copies of the PV data from source to target. If not selected, then it will perform an indirect copy, where the data is pushed to s3 object storage location via velero/restic. Then pulled from the s3 object store to the PV. The direct method is faster, since it performs a direct copy, as long as the two clusters have network connectivity to one another.

Create a migration plan

1 General

2 Namespaces

3 Persistent volumes

4 Copy options

5 Migration options

6 Hooks

Migration options

Images

Direct image migration Unavailable

☐ Use direct image migration

Persistent volumes

Direct PV migration Available

☒ Use direct PV migration for filesystem copies

Next

Back

Cancel

**Hooks:** For CPD migration, no hooks are used. Click "Next".

Create a migration plan

1 General

2 Namespaces

3 Persistent volumes

4 Copy options

5 Migration options

6 Hooks

Hooks

Hooks are commands that can be run at various steps in the migration process. They are defined in a container image or an Ansible playbook and can be run on either the source or target cluster.

Add hook

Name	Image	Destination	Type	Migration step
<div><div></div><div>No hooks have been added to this migration plan.</div></div>				

Next

Back

Cancel

**Migration plan validation:**

The migration plan validation might display the following known warning, which indicates that it cannot calculate how much capacity for each PVC has been used in the current state. So just make sure that you have not exceeded the capacity of any of the PVCs, otherwise the data transfer will fail. But as long as the applications honor the PVC limit size, this will not be a problem.

Failed to compute PV resizing data for the following volumes. PV resizing will be disabled for these volumes and the migration may fail if the volumes are full or their requested and actual capacities differ in the source cluster.

## Migration Plan PVC storage class validation

After the migration plan is created, you can see the migplan CR:

Example command:

```
$ oc get migplan -n openshift-migration
NAME                                READY    SOURCE    TARGET    STORAGE
AGE
cp4d-465-wvyb-to-9c9f             True     host      target-wvyb  ibmcos-mtc-todd-
repo    34m
$
```

Run the following command against the migration plan to see if the PVCs have the right matching of storage class to access mode: Command:

```
oc get migplan -n openshift-migration <migplan-name> -o json | jq -r
'.spec.persistentVolumes[] | "[PVC:" + .pvc.name + "]"Access mode:" +
.pvc.accessModes[0] + "]"New storage class:" + .selection.storageClass +
"]"'
```

Compare the output one by one to validate it is as expected.

Below are grep commands to help validate quickly:

- Validate that all RWX have cephfs storage class (you want no grep matches):

```
oc get migplan -n openshift-migration <migplan-name> -o json | jq -r
'.spec.persistentVolumes[] | "[PVC:" + .pvc.name + "]"Access mode:" +
.pvc.accessModes[0] + "]"New storage class:" + .selection.storageClass +
"]"' |grep ReadWriteMany |grep -v cephfs
```

- Validate that all RWO have ceph-rbd storage class (you want no grep matches):

```
oc get migplan -n openshift-migration <migplan-name> -o json | jq -r
'.spec.persistentVolumes[] | "[PVC:" + .pvc.name + "]"Access mode:" +
.pvc.accessModes[0] + "]"New storage class:" + .selection.storageClass +
"]"' |grep ReadWriteOnce |grep -v ceph-rbd
```

- Example output:

```
$ oc get migplan -n openshift-migration <migplan-name> -o json | jq -r
'.spec.persistentVolumes[] | "[PVC:" + .pvc.name + "][Access mode:" +
.pvc.accessModes[0] + "][New storage class:" + .selection.storageClass +
"]"' |grep ReadWriteMany |grep -v cephfs
$

$ oc get migplan -n openshift-migration <migplan-name> -o json | jq -r
'.spec.persistentVolumes[] | "[PVC:" + .pvc.name + "][Access mode:" +
.pvc.accessModes[0] + "][New storage class:" + .selection.storageClass +
"]"' |grep ReadWriteOnce |grep -v ceph-rbd
$
```

If the target storage classes for any of the PVCs do not match what you intended, then you will need to edit the Migration Plan and adjust the target storage classes. This can be accomplished by one of these methods:

- Go to the MTC console, Migration plans. Select the Migration Plan and edit.
- Edit the Migration Plan from cli: `oc edit migplan -n openshift-migration <migplan-name>`

## Shutdown Cloud Pak for Data instance on source cluster

---

### Set oadp prehooks pre-reqs

Reference: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=backup-prerequisite-tasks>

DB2 instances:

Obtain all the db2uclusters in the namespace

```
oc -n <cpd-namespace> get db2ucluster
```

For each of the db2ucluster displayed in the command above execute the following command:

```
oc -n <cpd-instance> label db2ucluster <DB2UCLUSTER> db2u/cpdbr=db2u --
overwrite
```

Verify each db2ucluster has the label `db2u/cpdbr=db2u`:

```
oc -n <cpd-instance> get db2ucluster --show-labels
```

### Scale down CPD services by running backup prehooks

Example command:

```
cpd-cli oadp backup prehooks --include-namespaces <cpd-namespace> --log-level=debug --verbose
```

Check the state of the deployments and statefulsets in the CPD instance namespace by running the following commands:

```
oc get deploy -n <cpd-namespace>
oc get sts -n <cpd-namespace>
oc get pod -n <cpd-namespace> --no-headers |grep -v Completed |wc -l
```

Run backup prehooks:

```
cpd-cli oadp backup prehooks --include-namespaces <cpd-namespace> --log-level=debug --verbose
```

Verify successful completion of backup prehooks by ensuring output for command above includes these messages:

```
configmap pre-backup hooks completed
backup prehooks command completed
```

After running the backup prehooks, check for the remaining deployments/statefulsets/pods running:

Example:

```
$ oc get deploy -n <cpd-namespace> |grep -v "0/0"
NAME                                READY    UP-TO-DATE
AVAILABLE    AGE
ibm-nginx    2/2        2            2
16h
manta-admin-gui    1/1        1            1
13h
manta-artemis    1/1        1            1
13h
manta-configuration-service    1/1        1            1
13h
manta-dataflow    1/1        1            1
13h
manta-flow-agent    1/1        1            1
13h
$
$ oc get sts -n <cpd-namespace> |grep -v "0/0"
```



```

NAME                                READY   AGE
c-db2oltp-iis-db2u                 1/1     13h
c-db2oltp-wkc-db2u                 1/1     13h
$
$ oc get pod -n <cpd-namespace> |grep -v Completed
NAME                                READY   STATUS
RESTARTS      AGE
c-db2oltp-iis-db2u-0                 1/1     Running
0                13h
c-db2oltp-wkc-db2u-0                 1/1     Running
0                13h
ibm-nginx-69dd4dc867-9vrdz           1/1     Running
0                16h
ibm-nginx-69dd4dc867-lc7pf           1/1     Running
0                16h
manta-admin-gui-5bcd558778-bq9mz     1/1     Running
0                13h
manta-artemis-8f4b5d55f-xmjfk        1/1     Running
0                13h
manta-configuration-service-5d87768b87-49bsw 1/1     Running
0                13h
manta-dataflow-56577f6f75-kxtss      1/1     Running
0                13h
manta-flow-agent-5dfcb4864b-rfd2n    1/1     Running
0                13h
wkc-foundationdb-cluster-cluster-controller-1 2/2     Running
0                14h
wkc-foundationdb-cluster-log-1        2/2     Running
0                14h
wkc-foundationdb-cluster-log-2        2/2     Running
0                14h
wkc-foundationdb-cluster-storage-1    2/2     Running
0                14h
wkc-foundationdb-cluster-storage-2    2/2     Running
1 (13h ago) 14h
wkc-foundationdb-cluster-storage-3    2/2     Running
0                14h
$

```

Scale down remaining services and resources running in CPD namespace.

For the remaining pods running in the CPD instance namespace, they will have to be shutdown manually. The following is a list of known services/deployments/statefulsets that have to be shutdown manually, and the process to do so.

#### nginx:

```

oc get deploy/ibm-nginx -n <cpd-namespace> #retain this output for the
resume
oc scale --replicas=0 deploy/ibm-nginx -n <cpd-namespace>

```

**db2uclusters statefulsets:**

```
oc get sts -n <cpd-namespace> |grep -E "NAME|db2" #retain this output for
the resume
oc scale --replicas=0 sts/c-db2oltp-iis-db2u -n <cpd-namespace>
oc scale --replicas=0 sts/c-db2oltp-wkc-db2u -n <cpd-namespace>

## Get sts for your db2wh instance
oc get sts -licpdsupport/addOnId=db2wh -n <cpd-namespace> #retain this
output for the resume, you will see two sts (one for db2u and one for
etcd)
oc scale --replicas=0 sts/<db2wh-db2u-sts> -n <cpd-namespace>
oc scale --replicas=0 sts/<db2wh-etcd-sts> -n <cpd-namespace>
```

Note: There may be additional db2u related statefulsets that need to be scaled down. Scale them down similar to the above commands.

**Foundationdb (FDB)**

```
oc patch -p '{"spec":{"shutdown":"true"}}' --type=merge fdbcluster wkc-
foundationdb-cluster -n <cpd-namespace>
```

**MANTA** (IBM internal reference: <https://github.ibm.com/manta/manta-adl-operator#scaling>)

```
oc patch -p '{"spec":{"replicas":0}}' --type=merge mantaflow mantaflow-wkc
-n <cpd-namespace>
```

**Openscale**

Note: Even though the openscale deployments and statefulsets are shutdown during the backup prehooks, the service itself is not shutdown and not in maintenance mode. That means that it could be reconciled and started up again when not desired on the target cluster after the cpd operators are restored. To prevent this, we need to manually shut down the openscale service on the source cluster.

```
oc patch -p '{"spec":{"shutdown":"true"}}' --type=merge woservice
aiopenscale -n <cpd-namespace>
oc patch -p '{"spec":{"ignoreForMaintenance": true}}' --type=merge
woservice aiopenscale -n <cpd-namespace>
```

**Datastage deploy and statefulset**

```
## Record this output for the resume operations:
oc get deploy,sts -n <cpd-namespace>|grep ds-px-default
```

```
oc scale --replicas=0 deploy/ds-px-default-ibm-datastage-px-runtime -n
<cpd-namespace>
oc scale --replicas=0 sts/ds-px-default-ibm-datastage-px-compute -n <cpd-
namespace>
```

### Example Output for scale down and shutdown commands above:

#### nginx

```
$ oc get deploy ibm-nginx -n cpd
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
ibm-nginx     2/2      2              2             16h
$ oc scale --replicas=0 deploy/ibm-nginx
deployment.apps/ibm-nginx scaled
$ oc get deploy ibm-nginx -n cpd
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
ibm-nginx     0/0      0              0             16h
$ oc get pod -n cpd |grep -v Completed |grep -E "NAME|nginx"
NAME                                                  READY    STATUS
RESTARTS      AGE
$
```

#### db2ucluster statefulsets

```
$ oc get db2ucluster -n cpd
NAME                STATE    MAINTENANCESTATE    AGE
db2oltp-iis         Ready    InMaintenance        2d17h
db2oltp-wkc         Ready    InMaintenance        8d
db2wh-1686277435040922 Ready    InMaintenance        13h

$ oc get sts -n cpd |grep -v "0/0"
NAME                READY    AGE
c-db2oltp-iis-db2u  1/1     2d17h
c-db2oltp-wkc-db2u  1/1     8d
c-db2wh-1686277435040922-db2u 1/1     13h
c-db2wh-1686277435040922-etcd 1/1     13h
ds-px-default-ibm-datastage-px-compute 2/2     8d

$ oc scale --replicas=0 sts/c-db2oltp-iis-db2u -n cpd
statefulset.apps/c-db2oltp-iis-db2u scaled
$ oc scale --replicas=0 sts/c-db2oltp-wkc-db2u -n cpd
statefulset.apps/c-db2oltp-wkc-db2u scaled
$ oc scale --replicas=0 sts/c-db2wh-1686277435040922-db2u -n cpd
statefulset.apps/c-db2wh-1686277435040922-db2u scaled
$ oc scale --replicas=0 sts/c-db2wh-1686277435040922-etcd -n cpd
statefulset.apps/c-db2wh-1686277435040922-etcd scaled
```

```
$ oc get sts -n cpd |grep -E "NAME|db2u"
NAME                                READY    AGE
c-db2oltp-iis-db2u                 0/0      14h
c-db2oltp-wkc-db2u                 0/0      14h
c-db2wh-1686277435040922-db2u     0/0      13h
c-db2wh-1686277435040922-etcd     0/0      13h

$ oc get pod -n cpd |grep -v Completed |grep -E "NAME|db2"
NAME                                READY    STATUS
RESTARTS    AGE
$
```

## Foundationdb shutdown

```
$ oc describe fdbcluster wkc-foundationdb-cluster -n cpd |grep "
foundationdb.opencontent.ibm.com/backup-trigger"
foundationdb.opencontent.ibm.com/backup-trigger: pre-backup
$
$ oc patch -p '{"spec":{"shutdown":"true"}}' --type=merge fdbcluster wkc-
foundationdb-cluster -n cpd
fdbcluster.foundationdb.opencontent.ibm.com/wkc-foundationdb-cluster
patched
$
$ oc get fdbcluster wkc-foundationdb-cluster -n cpd -o yaml |grep shutdown
shutdown: "true"
shutdownStatus: ' '
$

$ oc get pod -n cpd |grep -E "NAME|fdb|found"
NAME                                READY    STATUS
RESTARTS    AGE
wkc-foundationdb-cluster-cluster-controller-1    2/2
Terminating    0      14h
wkc-foundationdb-cluster-log-1                  2/2
Terminating    0      14h
wkc-foundationdb-cluster-log-2                  2/2
Terminating    0      14h
wkc-foundationdb-cluster-storage-1              0/2
Terminating    0      14h
wkc-foundationdb-cluster-storage-2              2/2
Terminating    1 (14h ago)  14h
wkc-foundationdb-cluster-storage-3              2/2
Terminating    0      14h
$
$ oc get fdbcluster wkc-foundationdb-cluster -n cpd -o yaml |grep shutdown
shutdown: "true"
shutdownStatus: shutdown
$
$ oc get pod -n cpd |grep -E "NAME|fdb|found"
NAME                                READY    STATUS
RESTARTS    AGE
$
```

MANTA scaledown replica

```
$ oc get pod -n cpd |grep -v Completed
NAME                                READY   STATUS
RESTARTS   AGE
manta-admin-gui-5bcd558778-bq9mz    1/1     Running
0           13h
manta-artemis-8f4b5d55f-xmjfk       1/1     Running
0           13h
manta-configuration-service-5d87768b87-49bsw 1/1     Running
0           13h
manta-dataflow-56577f6f75-kxtss     1/1     Running
0           13h
manta-flow-agent-5dfcb4864b-rfd2n    1/1     Running
0           13h
$
$ oc get mantaflow -n cpd
NAME          AGE
mantaflow-wkc 13h
$ oc get mantaflow -n cpd mantaflow-wkc -o yaml |grep replicas
replicas: 1
$
$ oc patch -p '{"spec":{"replicas":0}}' --type=merge mantaflow mantaflow-wkc -n cpd
mantaflow.adl.getmanta.com/mantaflow-wkc patched
$ oc get mantaflow mantaflow-wkc -o yaml |grep replicas
replicas: 0
$
$ oc get pod -n cpd |grep -E "NAME|manta"
NAME                                READY   STATUS
RESTARTS   AGE
$
```

Openscale shutdown:

Shutdown openscale:

```
$ oc -n cpd get woservice
NAME                                TYPE          STORAGE  SCALECONFIG
PHASE    RECONCILED  STATUS
aiopenscale          service          small
Ready    4.6.5       Completed
openscale-defaultinstance serviceInstance  small
Ready    4.6.5       Completed
$
$ oc get woservice aiopenscale -o json |jq .spec.shutdown
"false"
$
```

```
$ oc patch -p '{"spec":{"shutdown":"true"}}' --type=merge woservice
aiopenscale -n cpd
woservice.wos.cpd.ibm.com/aiopenscale patched
$ oc -n cpd get woservice aiopenscale -o json |jq -y .spec.shutdown
'true'
$
```

Place it in maintenance mode:

```
$ oc patch -p '{"spec":{"ignoreForMaintenance": true}}' --type=merge
woservice aiopenscale -n cpd
$ oc get woservice
NAME                                TYPE                                STORAGE  SCALECONFIG
PHASE      RECONCILED  STATUS
aiopenscale                service                                small
InMaintenance    4.6.5      InMaintenance
openscale-defaultinstance serviceInstance                        small
shutdown          4.6.5      shutdown
$ oc get woservice aiopenscale -o yaml
apiVersion: wos.cpd.ibm.com/v1
kind: WOService
metadata:
  creationTimestamp: "2023-05-19T05:34:44Z"
  finalizers:
  - wos.cpd.ibm.com/finalizer
  generation: 5
  name: aiopenscale
  namespace: cpdinstance
  resourceVersion: "13630612"
  uid: 18596e99-8ffc-47f4-9068-a0d2d3fd8542
spec:
  acceptRollback: false
  blockStorageClass: cpd-storage
  fileStorageClass: cpd-storage
  ignoreForMaintenance: true
  license:
    accept: true
    license: Enterprise
  scaleConfig: small
  shutdown: "true"
  type: service
  version: 4.6.5
status:
  conditions:
  - ansibleResult:
      changed: 1
      completion: 2023-05-19T18:49:06.575661
      failures: 0
      ok: 2
      skipped: 0
      lastTransitionTime: "2023-05-19T18:49:01Z"
      message: Awaiting next reconciliation
```

```

    reason: Successful
    status: "True"
    type: Running
  - lastTransitionTime: "2023-05-19T18:49:06Z"
    message: Last reconciliation succeeded
    reason: Successful
    status: "True"
    type: Successful
  - lastTransitionTime: "2023-05-19T18:49:01Z"
    message: ""
    reason: ""
    status: "False"
    type: Failure
phase: InMaintenance
versions:
  reconciled: 4.6.5
wosBuildNumber: "105"
wosStatus: InMaintenance
$

```

## Datastage

```

$ oc scale --replicas=0 deploy/ds-px-default-ibm-datastage-px-runtime -n
cpd
deploy.apps/ds-px-default-ibm-datastage-px-runtime scaled
$ oc scale --replicas=0 sts/ds-px-default-ibm-datastage-px-compute -n cpd
statefulset.apps/ds-px-default-ibm-datastage-px-compute scaled

```

## Validate all CPD services and pods are now shutdown:

Ensure all pods in the CPD instance namespace are now shutdown:

```

$ oc get pod -n <cpd-namespace> |grep -v Completed
NAME                                     READY   STATUS
RESTARTS   AGE
$

```

Ensure all CPD services are in maintenance mode:

```

for cpdcrd in $(oc get crd |grep cpd.ibm.com |awk '{print $1}');do echo --
- $cpdcrd ---;oc get $cpdcrd -o name |xargs -I{} sh -c "oc get {} -o yaml
|grep -i maintenance";done

```

Example:

```
$ for cpdcrd in $(oc get crd |grep cpd.ibm.com |awk '{print $1}');do echo
--- $cpdcrd ---;oc get $cpdcrd -o name |xargs -I{} sh -c "oc get {} -o
yaml |grep -i maintenance";done
--- analyticsengines.ae.cpd.ibm.com ---
  ignoreForMaintenance: true
  analyticsengineStatus: InMaintenance
--- ccs.ccs.cpd.ibm.com ---
  ignoreForMaintenance: true
  ccsStatus: InMaintenance
--- datarefinery.datarefinery.cpd.ibm.com ---
  ignoreForMaintenance: true
  datarefineryStatus: InMaintenance
--- datastages.ds.cpd.ibm.com ---
--- db2aaserviceservices.databases.cpd.ibm.com ---
  ignoreForMaintenance: true
  db2aaserviceStatus: InMaintenance
--- db2oltpservices.databases.cpd.ibm.com ---
  ignoreForMaintenance: true
  db2oltpStatus: InMaintenance
--- db2whservices.databases.cpd.ibm.com ---
  ignoreForMaintenance: true
  db2whStatus: InMaintenance
--- ibmcpds.cpd.ibm.com ---
--- iis.iis.cpd.ibm.com ---
  ignoreForMaintenance: true
  iisStatus: InMaintenance
--- notebookruntimes.ws.cpd.ibm.com ---
  ignoreForMaintenance: true
  runtimeStatus: InMaintenance
--- pxruntimes.ds.cpd.ibm.com ---
--- ug.wkc.cpd.ibm.com ---
  ignoreForMaintenance: true
  ugStatus: InMaintenance
--- wkc.wkc.cpd.ibm.com ---
  ignoreForMaintenance: true
  wkcStatus: InMaintenance
--- wmlbases.wml.cpd.ibm.com ---
  ignoreForMaintenance: true
  wmlStatus: InMaintenance
--- woservices.wos.cpd.ibm.com ---
  ignoreForMaintenance: true
  phase: InMaintenance
  wosStatus: InMaintenance
  ignoreForMaintenance: false
--- ws.ws.cpd.ibm.com ---
  ignoreForMaintenance: true
  wsStatus: InMaintenance
--- zenextensions.zen.cpd.ibm.com ---
--- zenservices.zen.cpd.ibm.com ---
  ignoreForMaintenance: true
  zenStatus: InMaintenance
```



## MTC migration Stage to target cluster

The migration Stage will create the corresponding namespace on the target cluster. Then it will create the PVs and PVCs and copy the current data to the target cluster PVs, using whichever copy options were selected in the migration plan.

Note: The "Stage" process is concurrent and can be performed while the CPD services are running. No application resources will be scaled down or shutdown in this phase.

**On the target cluster, verify that the cpd namespace to be migrated does not exist yet:** Example command:

```
$ oc get ns |grep -v open
```

NAME	STATUS	AGE
default	Active	13d
ibm-common-services	Active	5d21h
kube-node-lease	Active	13d
kube-public	Active	13d
kube-system	Active	13d
oadp-operator	Active	5d22h
turbo	Active	13d

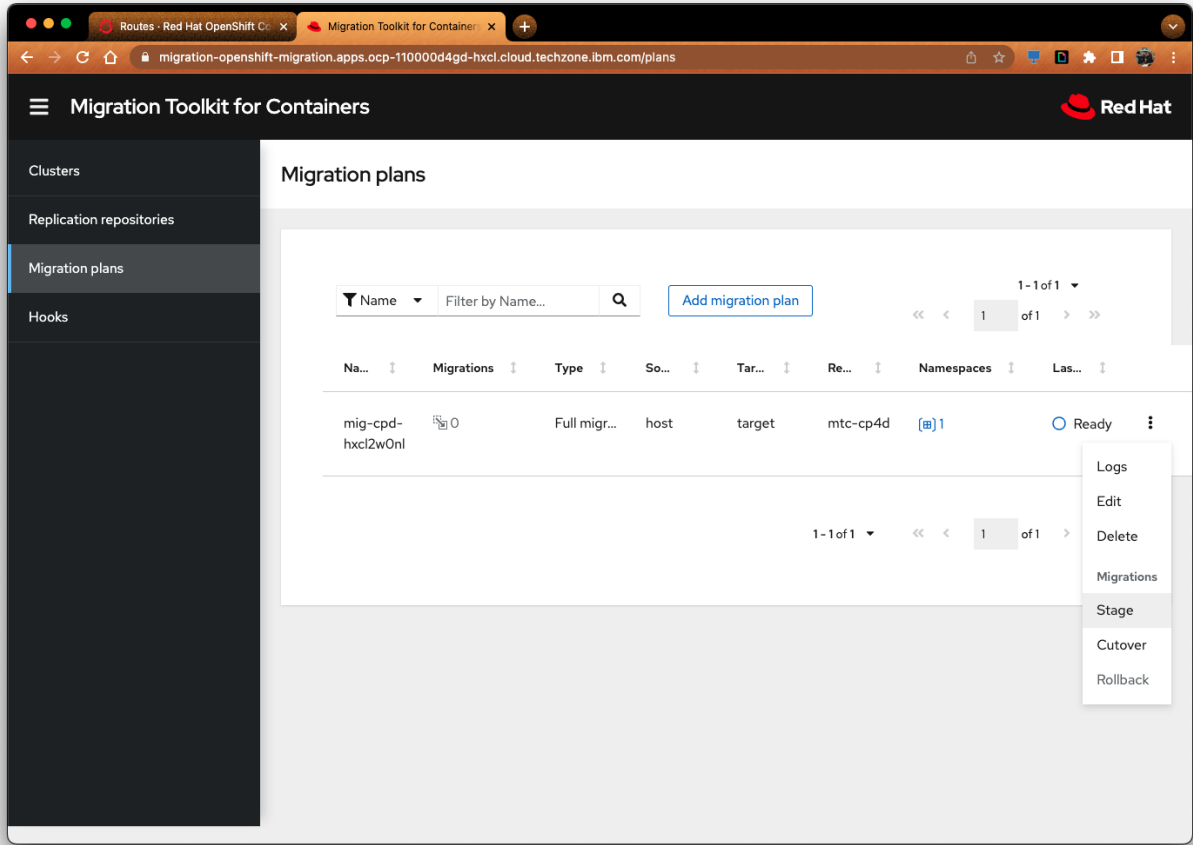
**Verify that the PVs do not exist on the target cluster:**

```
$ oc get pv
```

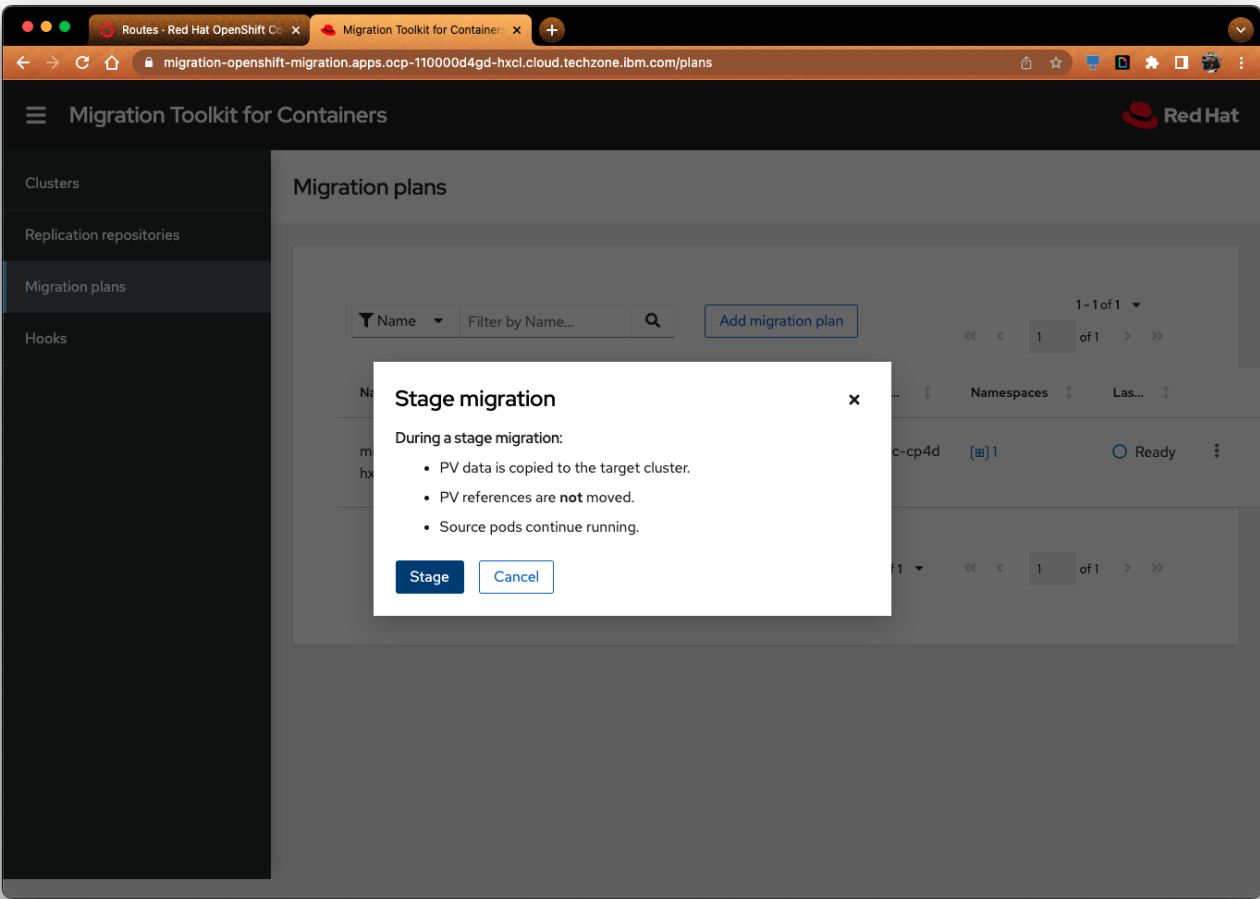
NAME	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	CAPACITY	ACCESS MODES
local-pv-18e3632b	Delete	Bound	openshift-storage/ocs-deviceset-2-data-0s1s6n	localblock	700Gi	RWO
local-pv-1f601e15	Delete	Bound	openshift-storage/ocs-deviceset-1-data-0hxdkm	localblock	700Gi	RWO
local-pv-d647fb26	Delete	Bound	openshift-storage/ocs-deviceset-0-data-099z84	localblock	700Gi	RWO
pvc-22cad882-f36a-41d1-8efb-e954859e86d8	Delete	Bound	busybox-6y5e/busybox-pvc	ocs-storagecluster-cephfs-alt-cluster	1Gi	RWO
pvc-8c4d2bab-3815-42f8-bd23-004db1913711	Delete	Bound	openshift-storage/db-noobaa-db-pg-0	ocs-storagecluster-ceph-rbd	50Gi	RWO
pvc-bceec5e0-75bd-4e35-8652-b106cbf00cc1	Delete	Bound	busybox-migrate/busybox-cephfs-pvc	ocs-storagecluster-cephfs-alt-cluster	5Gi	RWX
pvc-dd1054a8-f7a9-4500-82ef-4172eddb0fe	Delete	Bound	busybox-6y5e/busybox-pvc-new	ocs-storagecluster-cephfs-alt-cluster	1Gi	RWO

Perform Migration Stage:

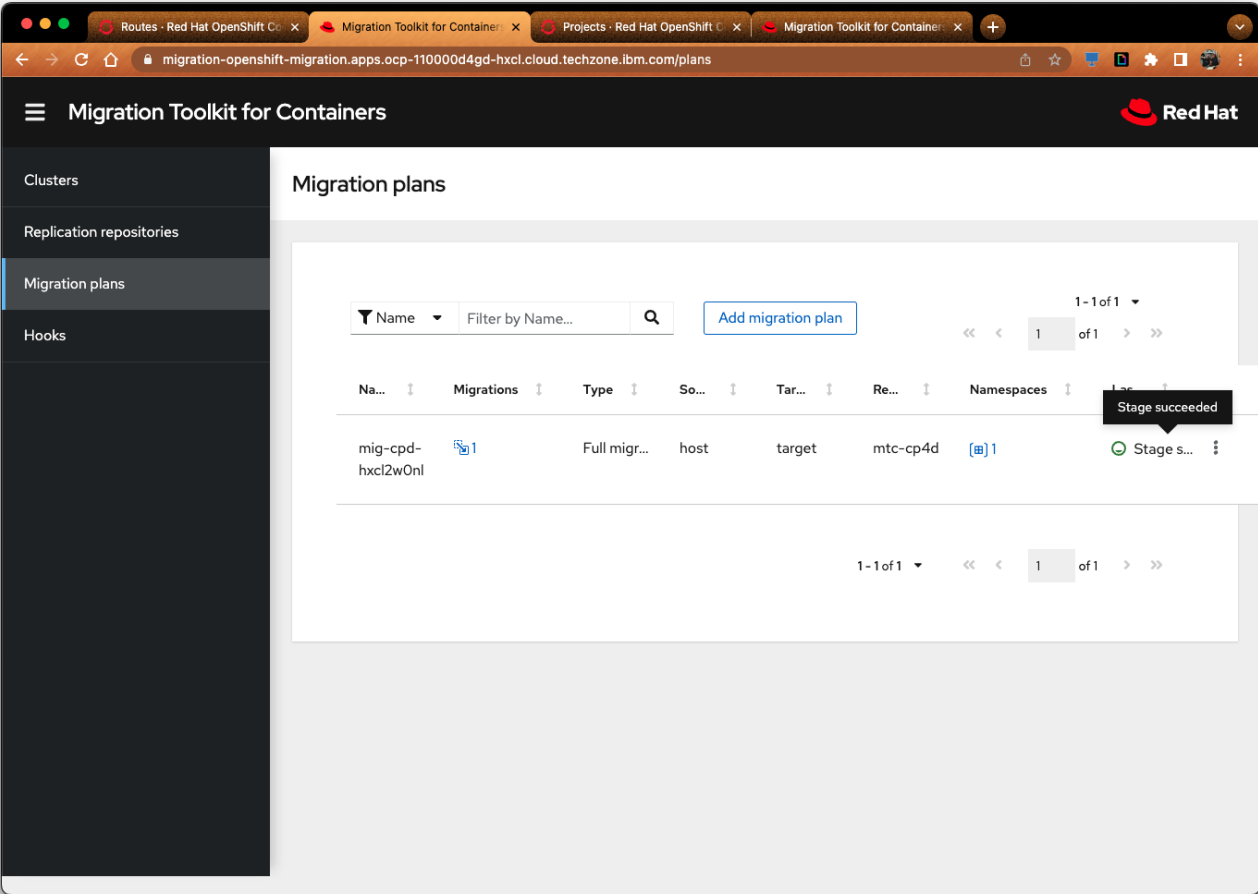
- From MTC console on source cluster, go to "Migration plans". For the CPD migration plan, go to the three dots on the right side and select "Stage". This will create the PVs and PVCs in the cpd namespace on the target cluster.

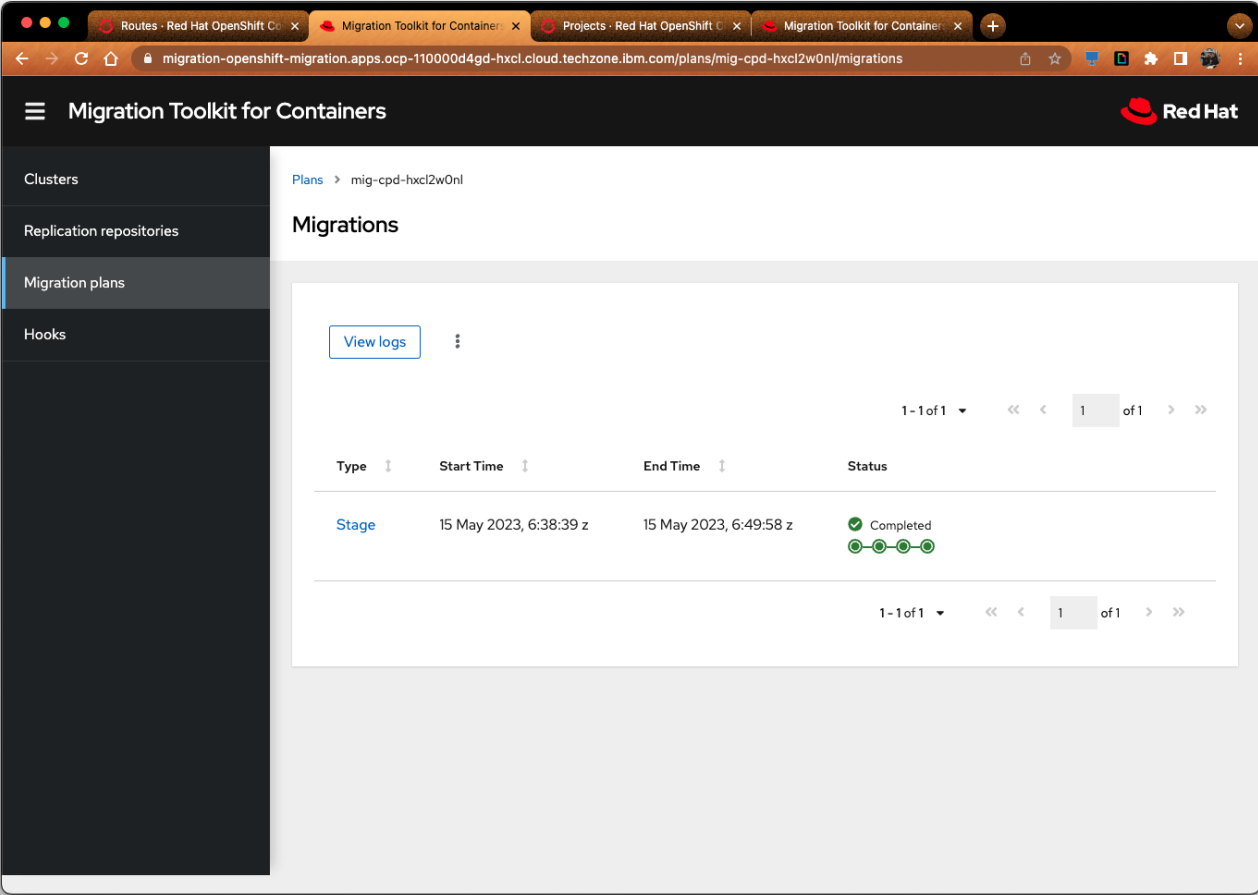


Click **Stage**



Ensure Stage completes successfully with no warnings:





During the Stage process, the **cpd** namespace, PVs and PVCs will be created and the **rsync-server** pod will be started.

Verify Cloud Pak for Data namespace was created.

```
oc get ns <cpd-namespace>
```

Note: During the Stage process, you may see an error indicating "Migration may be stuck." creating the rsync pods. Be patient in this phase, as the rsync pods will usually create successfully and the migration activity will complete.

Migration Toolkit for Containers

Clusters

Replication repositories

Migration plans

Hooks

Plans > cp4d-465-wvyb-to-9c9f > Cutover - 21 May 2023, 20:41:16

Migration details

⚠ Migration may be stuck.

Repeated errors occurred when attempting to create one or more Rsync pods in the source cluster. Please check controller logs for details.

Step	Elapsed time	Status	
Prepare	a few seconds	Complete	
Backup	a minute	Complete	<div>Backup openshift-migration/migration-d4dc0-initial-8grtv</div> <div>100% ✓</div>
StageBackup	a few seconds	Complete	

If the Stage process fails, check velero status of the failed velero "migration" job. Example: `velero describe restore migration-xxx-final-xxx -n openshift-migration`

## Target cluster migration steps

Verify PVs are created on target cluster and verify PV total matches original total before migration.

```
oc get pv |grep <cpd-namespace>
oc get pv |grep <cpd-namespace> | wc -l
```

Verify PVCs are created on target cluster.

```
oc get pvc -n <cpd-namespace>
```

Verify `rsync-server` started by running the command below. Note: The rsync server on the target cluster will be removed when the PV rsync actions are complete.

```
$ oc get pods -n <cpd-namespace> | grep rsync
NAME          READY   STATUS    RESTARTS   AGE
rsync-server  2/2     Running   0           30s
```

Additionally, `rsync` pods will start and complete on the source cluster. Example output:

```
$ oc get pod -n <cpd-namespace> | grep rsync
rsync-22qck           0/2
Completed             0          35s
```

rsync-2hh4r			0/2
Completed	0	2m13s	
rsync-4f486			0/2
Completed	0	43s	
rsync-7hwrp			0/2
Completed	0	96s	
rsync-7xhh2			0/2
Completed	0	87s	
rsync-8wbwv			0/2
Completed	0	2m4s	
rsync-jb5gj			0/2
Completed	0	10s	
rsync-m7ks6			0/2
Completed	0	77s	
rsync-mnp7t			0/2
Completed	0	26s	
rsync-pdsbn			0/2
Completed	0	114s	
rsync-q9lcq			0/2
ContainerCreating	0	3s	
rsync-rmbmm			0/2
Completed	0	19s	
rsync-rn7zm			0/2
Completed	0	60s	
rsync-tcdjf			0/2
Completed	0	105s	
rsync-vbzfz			0/2
Completed	0	51s	
rsync-xmm5k			0/2
Completed	0	69s	

## MTC Migration Cutover to target cluster

Before cutover, validate again that no pods are running in source cluster and target cluster. This is needed to ensure that everything is stopped on source cluster for the cutover, so when the cutover is performed, MTC does not attempt to resume any operations. The services will be resumed manually for CPD in later steps.

Source cluster example:

```
$ oc get pod -n <cpd-namespace> |grep -v Completed
NAME                                     READY   STATUS
RESTARTS   AGE
$
```

Target cluster example:

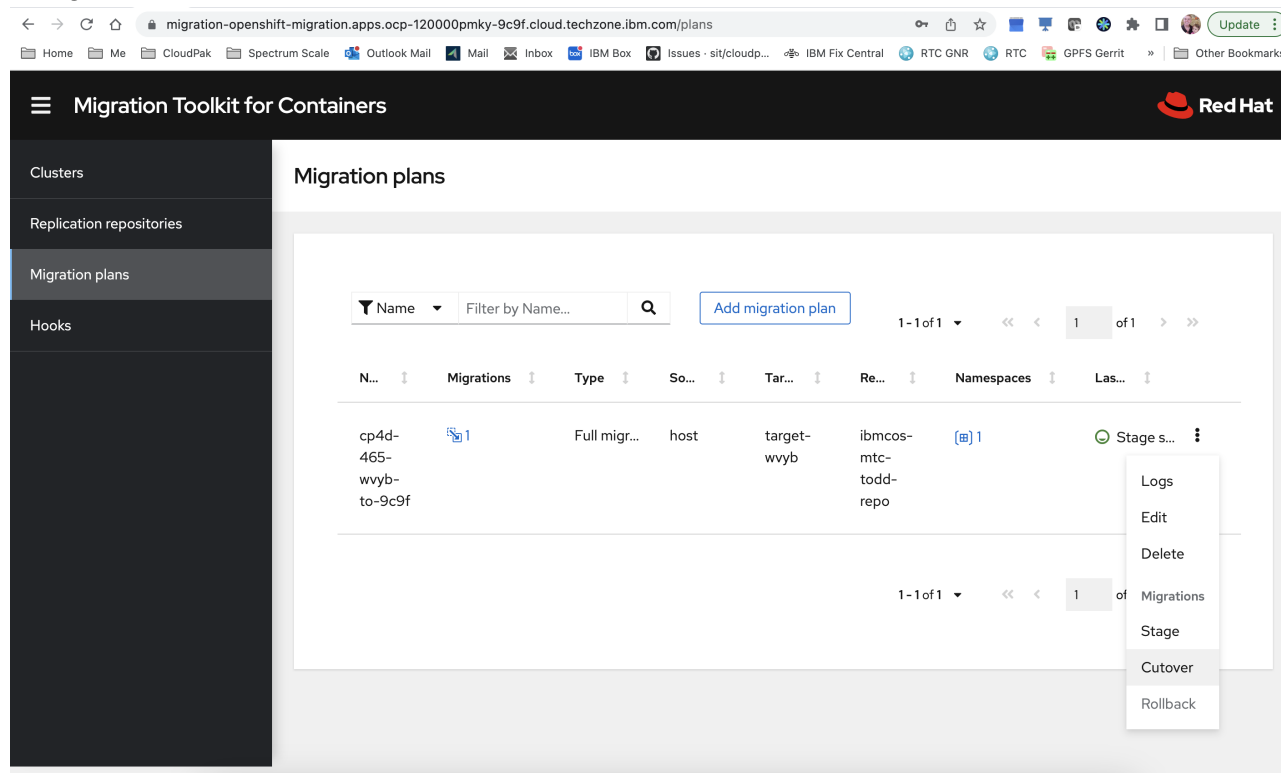
```
$ oc get pod -n <cpd-namespace>
No resources found in cpdinstance namespace.
$
```

## Perform the MTC Cutover from source cluster MTC UI:

Gather target cluster resources before cutover by running the following command:

```
$ NSPACE="<cpd-namespace>";oc get $(oc api-resources --namespaced=true --verbs=list -o name | awk '{printf "%s%s",sep,$0;sep=","}') --ignore-not-found -n ${NSPACE} |grep -v packagemanifest
```

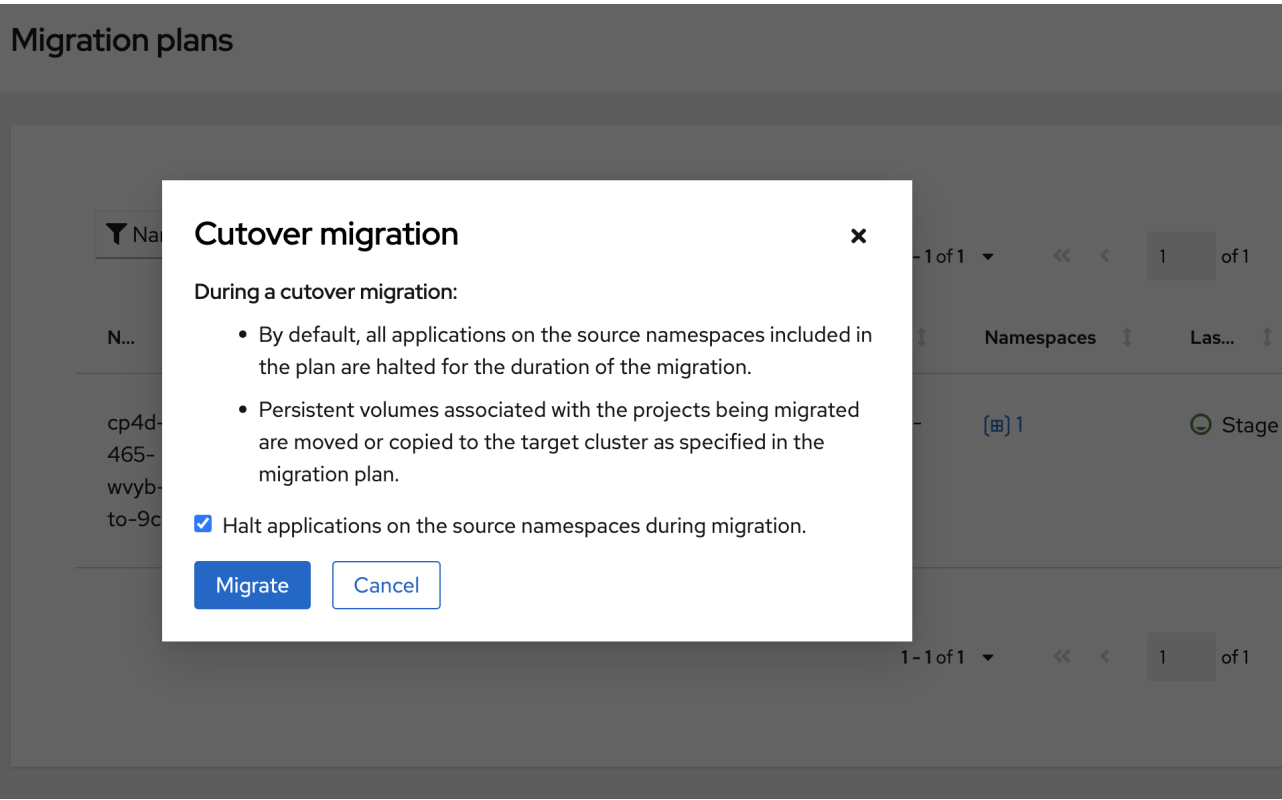
- From Migration Controller, go to "Migration plans". For the CPD migration plan, select the three dots on right hand side and select "Cutover".



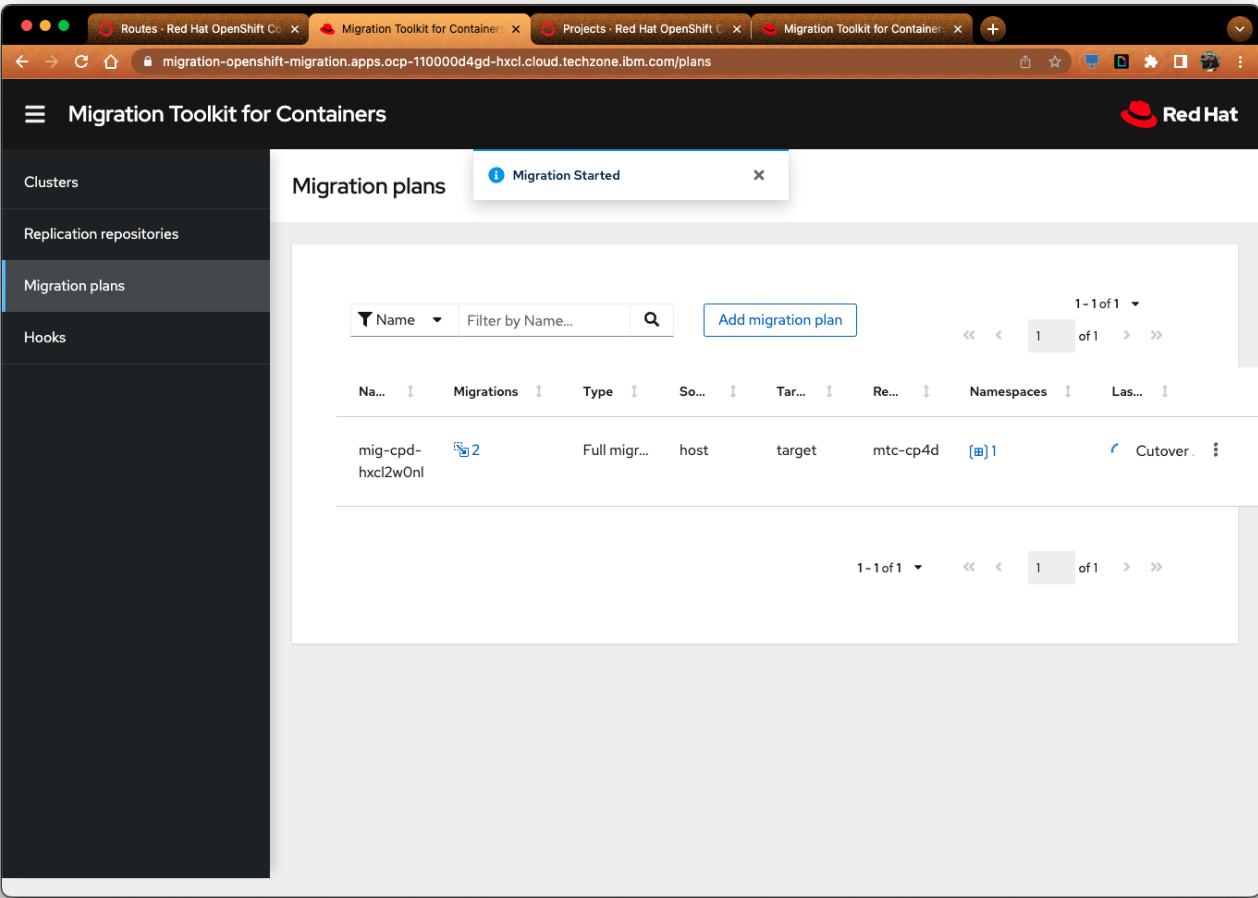
- A warning is displayed that applications will be halted. But this is not an issue, since we manually halted all applications prior to cutover. Make sure **Halt Applications...** option is selected and

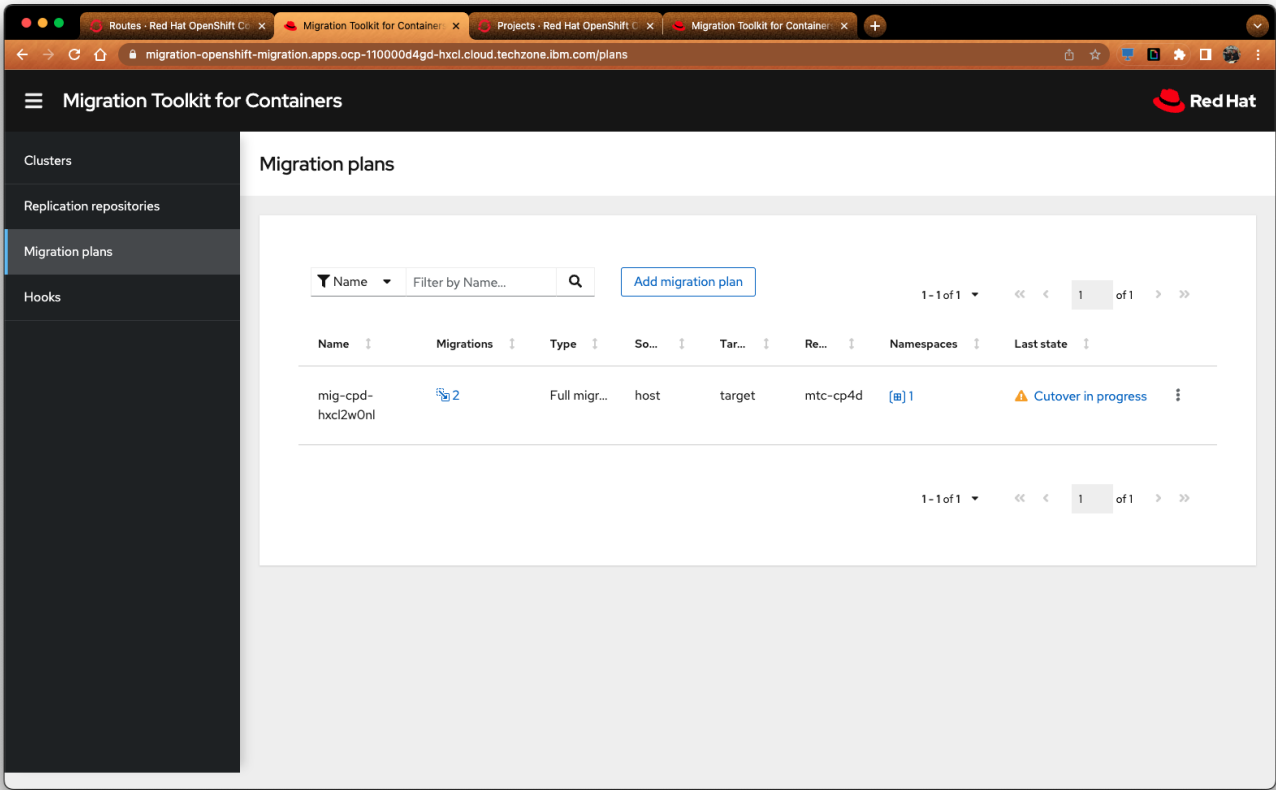


click **Migrate**

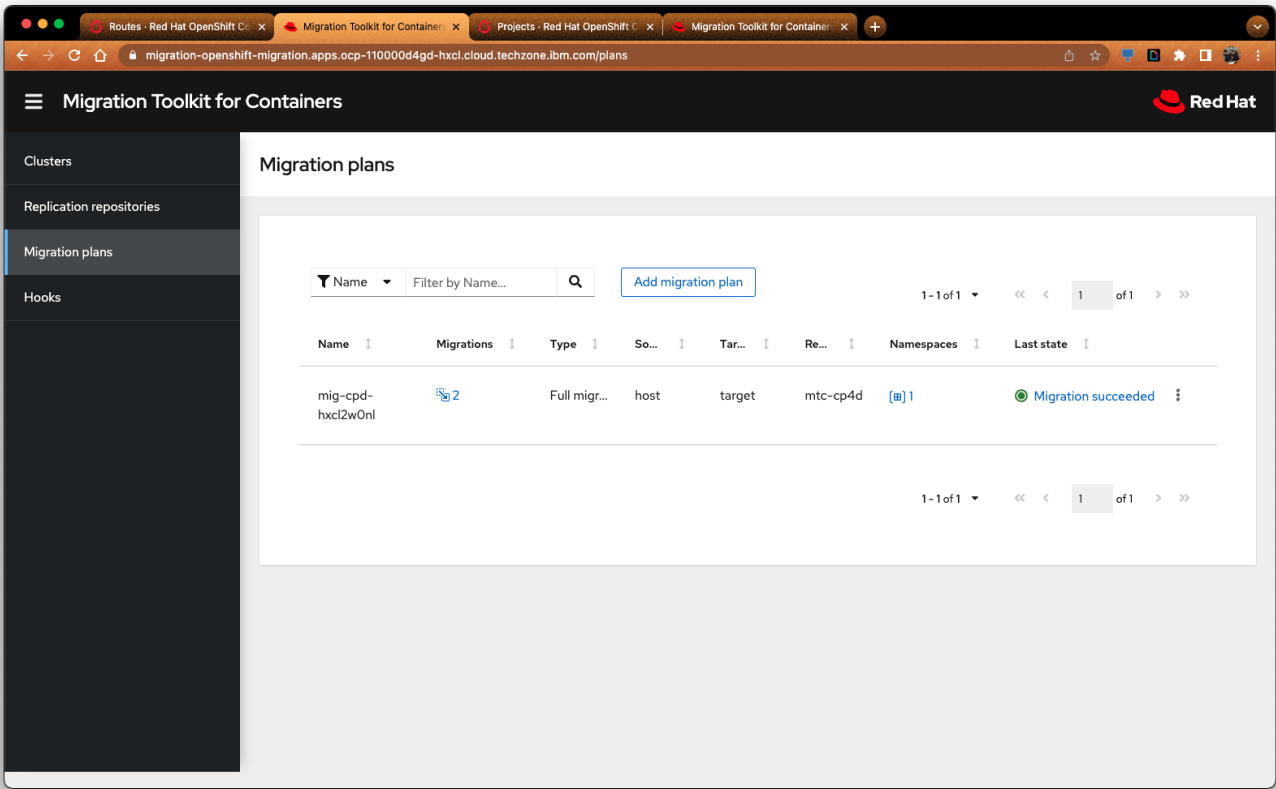


Migration cutover started





Migration completed:



During the migration cutover, an `rsync-server` pod will be created on the target cluster and `rsync` pods will start on the source cluster (similar to migration stage).

## rsync-server pod on target cluster

```
$ oc get pod -n <cpd-namespace> | grep rsync
```

NAME	READY	STATUS	RESTARTS	AGE
rsync-server	2/2	Running	0	8m16s

## rsync pods on source cluster

```
$ oc get pod -n <cpd-namespace> | grep rsync
```

rsync-4q5gc				0/2
Completed	0	89s		
rsync-4rbsk				0/2
Completed	0	32s		
rsync-6vs6g				0/2
Completed	0	3m10s		
rsync-9jtt7				0/2
Completed	0	2m9s		
rsync-bf6tl				0/2
Completed	0	4m14s		
rsync-c547g				0/2
Completed	0	3m52s		
rsync-cgwqh				0/2
Completed	0	109s		
rsync-dkhxx				0/2
Completed	0	2m49s		
rsync-hvd5m				0/2
Completed	0	7m45s		
rsync-lbxvl				0/2
Completed	0	4m36s		
rsync-lzdcn				0/2
Completed	0	6m32s		
rsync-m2g2p				0/2
Completed	0	8m2s		
rsync-m7wm4				0/2
Completed	0	6m56s		
rsync-ncbwl				0/2
Completed	0	3m31s		
rsync-nfn7f				0/2
Completed	0	4m59s		
rsync-pprkq				0/2
Completed	0	2m29s		
rsync-qzkrm				0/2
Completed	0	6m8s		
rsync-rdc54				0/2
Completed	0	51s		
rsync-rm9mh				0/2
Completed	0	5m45s		
rsync-sdzck				0/2
Completed	0	7m20s		
rsync-tw8j7				0/2

Completed	0	70s	
rsync-w54bc			0/2
Completed	0	13s	
rsync-xxxvz			0/2
Completed	0	5m22s	

Note: During the Cutover process, you may see an error indicating "Migration may be stuck." creating the rsync pods. Be patient in this phase, as the rsync pods will usually create successfully and the migration activity will complete.

If the Cutover process fails, check velero status of the failed velero "migration" job. Example: `velero describe restore migration-xxx-final-xxx -n openshift-migration`

## Migration Cutover Validation on Target cluster

Gather resources in CPD namespace:

```
$ oc get zenservice -n <cpd-namespace>
NAME      AGE
lite-cr   40m
$ oc get wkc -n <cpd-namespace>
NAME      VERSION  RECONCILED  STATUS  AGE
wkc-cr    4.6.5    true        Running 40m
$ oc get db2ucluser -n <cpd-namespace>
NAME      STATE  MAINTENANCESTATE  AGE
db2oltp-wkc  Active  MaintenanceState  40m
$ oc get ws
NAME      VERSION  RECONCILED  STATUS  AGE
ws-cr     6.5.0    true        Running 41m
$ oc get sts -n <cpd-namespace>
```

NAME	READY	AGE
aiopyscale-ibm-aio-etc	0/0	42m
aiopyscale-ibm-aio-kafka	0/0	42m
aiopyscale-ibm-aio-redis	0/0	42m
aiopyscale-ibm-aio-zookeeper	0/0	42m
c-db2oltp-wkc-db2u	0/0	42m
dsx-influxdb	0/0	42m
elasticsearch-master	0/0	42m
rabbitmq-ha	0/0	42m
redis-ha-server	0/0	42m
wdp-couchdb	0/0	42m
wml-cpd-etc	0/0	42m
wml-deployment-agent	0/0	42m
zen-metastoredb	0/0	42m

No pods found running in CPD namespace

```
$ oc get pod -n <cpd-namespace>
No resources found in cpd namespace.
```

Statefulsets will still point to the old storage class (Source storage class used in this example is called **cpd-storage**)

Example command:

```
$ oc get sts -n <cpd-namespace> c-db2oltp-wkc-db2u -o yaml | grep <cpd-storage-class>
storageClassName: cpd-storage
```

Migration cutover completed successfully.

## Patch and update new storage class references on target cluster

### Description of why patching is required:

MTC migration does not scour and patch all k8s resources for storage class references. It will only change the PV/PVC references, then copy over all of the existing k8s resources from the source namespace. However, any CR resources/etc that internally have references to storage classes and PVC claims for storage classes need to be adjusted to fit the new target cluster's storage classes.

Patch all kubernetes resources in the CPD namespace, by changing all references of old storage class(es) to new storage class(es).

- Find all resources with matches for the old storage class. Then strategically edit/change the corresponding k8s resource and change the referenced storage class.

- **Note:** You must make sure you pick the correct RWO or RWX storage class. Do not find/replace indiscriminately.

## Discover all resources in the cpd instance namespace with storage class references

If an oadp dump of the cpd namespace was collected in the earlier step "**Collect an oadp resources dump of the cpd namespace on source cluster**", that can be used to discover the resources that will need to be patched ahead of time. Do not attempt to collect an oadp resources dump if MTC was installed after cpd oadp was installed due to conflicts with oadp versions.

Alternatively, the resources can be queried live on either the source or target cluster.

Run the following command on the target cluster to look through all resources in the cpd namespace to find references to the old storage class, which will need to be patched/replaced with the new storage class(es). (For this command, replace `<cpd-namespace>` and `<source-storage-class>` with your values.)

```
NSPACE="<cpd-namespace>";SOURCE_SC="<source-storage-class>";for cr in $(oc
get $(oc api-resources --namespaced=true --verbs=list -o name | grep -vE
"packagemanifests|events|endpointslice|persistentvolumeclaim" | awk
'{printf "%s%s",sep,$0;sep=","}') --ignore-not-found -n ${NSPACE} -o
name);do cr_grep=$(oc get $cr -oyaml |grep -E "${SOURCE_SC}" |grep -v "
{"apiVersion\":"");if [[ -n "$cr_grep" ]];then echo $cr;fi;done
```

Example:

```
$ NSPACE="cpdinstance";SOURCE_SC="cpd-storage";for cr in $(oc get $(oc
api-resources --namespaced=true --verbs=list -o name | grep -vE
"packagemanifests|events|endpointslice|persistentvolumeclaim" | awk
'{printf "%s%s",sep,$0;sep=","}') --ignore-not-found -n ${NSPACE} -o
name);do cr_grep=$(oc get $cr -oyaml |grep -E "${SOURCE_SC}" |grep -v "
{"apiVersion\":"");if [[ -n "$cr_grep" ]];then echo $cr;fi;done
configmap/ibm-cpp-config
configmap/iis-db2u-config
configmap/wdp-profiling-iae-config
configmap/wkc-db2u-config
configmap/zen-lite-operation-configmap
mantaflow.adl.getmanta.com/mantaflow-wkc
analyticsengine.ae.cpd.ibm.com/analyticsengine-sample
statefulset.apps/aiopenscale-ibm-aio-etcd
statefulset.apps/aiopenscale-ibm-aio-kafka
statefulset.apps/aiopenscale-ibm-aio-zookeeper
statefulset.apps/c-db2oltp-iis-db2u
statefulset.apps/c-db2oltp-wkc-db2u
statefulset.apps/dsx-influxdb
statefulset.apps/elasticsearch-master
statefulset.apps/kafka
statefulset.apps/rabbitmq-ha
statefulset.apps/redis-ha-server
```

```

statefulset.apps/solr
statefulset.apps/wdp-couchdb
statefulset.apps/zen-metastoredb
statefulset.apps/zookeeper
foundationdbcluster.apps.foundationdb.org/wkc-foundationdb-cluster
ccs.ccs.cpd.ibm.com/ccs-cr
ibmcpd.cpd.ibm.com/ibmcpd-cr
db2oltpservice.databases.cpd.ibm.com/db2oltp-cr
datarefinery.datarefinery.cpd.ibm.com/datarefinery-sample
db2ucluster.db2u.databases.ibm.com/db2oltp-iis
db2ucluster.db2u.databases.ibm.com/db2oltp-wkc
formation.db2u.databases.ibm.com/db2oltp-iis
formation.db2u.databases.ibm.com/db2oltp-wkc
fdbcluster.foundationdb.opencontent.ibm.com/wkc-foundationdb-cluster
iis.iis.cpd.ibm.com/iis-cr
ug.wkc.cpd.ibm.com/ug-cr
wkc.wkc.cpd.ibm.com/wkc-cr
woservice.wos.cpd.ibm.com/aiopenscale
woservice.wos.cpd.ibm.com/openscale-defaultinstance
zenservice.zen.cpd.ibm.com/lite-cr
$

```

## Execute the cpd-migrate.sh script to patch the storage class references

### Prepare input.json file:

Before running the cpd-migrate.sh script, prepare the input.json file that will be fed into the cpd-migrate.sh script. This input.json file is expected to have matching json entries for every resource found in the discovery, from the previous step.

The creation of this input.json file was completed by the IBM team and it is called **cpd-migrate.input.json**. It has been provided for use with the **cpd-migrate.sh** script.

An example of the format of the input.json file:

```

{
  "analyticsengines.ae.cpd.ibm.com": {
    "analyticsengine-sample": [
      {
        "action": "patch",
        "type": "json-path",
        "sourcepath": ".spec.storageClass",
        "sourcevalue": "--source-file-storage-class",
        "targetvalue": "--target-file-storage-class"
      }
    ]
  },
  "statefulsets.apps": {
    "aiopenscale-ibm-aos-etcd": [
      {
        "action": "delete-create",

```

```

        "type": "json-path",
        "sourcepath":
".spec.volumeClaimTemplates[0].spec.storageClassName",
        "sourcevalue": "--source-block-storage-class",
        "targetvalue": "--target-block-storage-class"
    }
],
"zen-metastoredb": [
{
    "action": "delete-create",
    "type": "json-path",
    "sourcepath":
".spec.volumeClaimTemplates[0].spec.storageClassName",
    "sourcevalue": "--source-block-storage-class",
    "targetvalue": "--target-block-storage-class"
}
]
}
}

```

### Run the cpd-migrate.sh script in preview mode:

Run the cpd-migrate.sh in preview mode, which will output json files with the updated storage class references. The resulting json files should be inspected to validate that the new storage class references are accurate and contains valid content to be applied to patch the corresponding resource.

Command (assumes new storage classes are default ODF storage class names):

```

$ ./cpd-migrate.sh restore \
--input-file cpd-migrate.input.json \
--cpd-operand-namespace <cpd-namespace> \
--source-block-storage-class <source-cluster-nfs-sc> \
--target-block-storage-class ocs-storagecluster-ceph-rbd \
--source-file-storage-class <source-cluster-nfs-sc> \
--target-file-storage-class ocs-storagecluster-cephfs \
--output-directory /tmp/cpd-migrate \
--preview

```

### Example command:

```

./cpd-migrate.sh restore --input-file cpd-migrate.input.json --cpd-
operand-namespace cpdinstance --source-block-storage-class cpd-storage --
target-block-storage-class ocs-storagecluster-ceph-rbd --source-file-
storage-class cpd-storage --target-file-storage-class ocs-storagecluster-
cephfs --preview --output-directory ./preview

```

Example output-directory contents with resulting json files:



```
$ ls -l
cpdinstance.analyticsengines.ae.cpd.ibm.com.analyticsengine-sample.json
cpdinstance.ccs.ccs.cpd.ibm.com.ccs-cr.json
cpdinstance.configmaps.wdp-profiling-iae-config.json
cpdinstance.configmaps.wkc-db2u-config.json
cpdinstance.configmaps.zen-lite-operation-configmap.json
cpdinstance.datarefinery.datarefinery.cpd.ibm.com.datarefinery-sample.json
cpdinstance.db2uclusters.db2u.databases.ibm.com.db2oltp-iis.json
cpdinstance.db2uclusters.db2u.databases.ibm.com.db2oltp-wkc.json
cpdinstance.fdbclusters.foundationdb.opencontent.ibm.com.wkc-foundationdb-
cluster.json
cpdinstance.formations.db2u.databases.ibm.com.db2oltp-iis.json
cpdinstance.formations.db2u.databases.ibm.com.db2oltp-wkc.json
cpdinstance.foundationdbclusters.apps.foundationdb.org.wkc-foundationdb-
cluster.json
cpdinstance.ibmcpds.cpd.ibm.com.ibmcpd-cr.json
cpdinstance.iis.iis.cpd.ibm.com.iis-cr.json
cpdinstance.mantaflows.adl.getmanta.com.mantaflow-wkc.json
cpdinstance.statefulsets.apps.aiopenscale-ibm-aio-etcd.json
cpdinstance.statefulsets.apps.aiopenscale-ibm-aio-kafka.json
cpdinstance.statefulsets.apps.aiopenscale-ibm-aio-zookeeper.json
cpdinstance.statefulsets.apps.c-db2oltp-iis-db2u.json
cpdinstance.statefulsets.apps.c-db2oltp-wkc-db2u.json
cpdinstance.statefulsets.apps.dsx-influxdb.json
cpdinstance.statefulsets.apps.elasticsearch-master.json
cpdinstance.statefulsets.apps.kafka.json
cpdinstance.statefulsets.apps.rabbitmq-ha.json
cpdinstance.statefulsets.apps.redis-ha-server.json
cpdinstance.statefulsets.apps.solr.json
cpdinstance.statefulsets.apps.wdp-couchdb.json
cpdinstance.statefulsets.apps.zen-metastoredb.json
cpdinstance.statefulsets.apps.zookeeper.json
cpdinstance.ug.wkc.cpd.ibm.com.ug-cr.json
cpdinstance.wkc.wkc.cpd.ibm.com.wkc-cr.json
cpdinstance.woservices.wos.cpd.ibm.com.aiopenscale.json
cpdinstance.woservices.wos.cpd.ibm.com.openscale-defaultinstance.json
cpdinstance.zenservices.zen.cpd.ibm.com.lite-cr.json
```

### Run the cpd-migrate.sh script to perform the resource updates:

After running cpd-migrate.sh in preview mode, and inspecting and validating the resulting json files, run the cpd-migrate.sh script without the preview option. This will take action on the kubernetes resources in the actual cluster, updating the resources with storage class references in the cpd instance namespace.

**Note:** There will be a small subset of resources that will have actions of "ignore" or "manual". This means that the cpd-migrate.sh tool will not take action on those resources.

- For a resource with action of "ignore", that indicates that no action is required and the existing reference to the old storage class can remain.

Example of "ignore" action:

```
-----
Time: 2023-05-31T06:05:41.543+0000 level=info - patch-resource:
cpdinstance, configmaps, ibm-cpp-config, [{"action":"ignore","type":"json-
path","sourcepath":".data.storageclass.default","sourcevalue":"--source-
block-storage-class","targetvalue":""},{ "action":"ignore","type":"json-
path","sourcepath":".data.storageclass.list","sourcevalue":"--source-
block-storage-class","targetvalue":""}]
```

```
Time: 2023-05-31T06:05:43.186+0000 level=info - Migrate Action:
{"action":"ignore","type":"json-
path","sourcepath":".data.storageclass.default","sourcevalue":"--source-
block-storage-class","targetvalue":""}
```

```
Time: 2023-05-31T06:05:43.497+0000 level=warning - Ignore
.data.storageclass.default: "cpd-storage"
```

```
Time: 2023-05-31T06:05:43.500+0000 level=info - Migrate Action:
{"action":"ignore","type":"json-
path","sourcepath":".data.storageclass.list","sourcevalue":"--source-
block-storage-class","targetvalue":""}
```

```
Time: 2023-05-31T06:05:43.809+0000 level=warning - Ignore
.data.storageclass.list: "cpd-storage"
```

```
Time: 2023-05-31T06:05:43.812+0000 level=info - configmaps ibm-cpp-config
- No Change
-----
```

- For a resource with action of "manual", that indicates that the cpd-migrate.sh tool is not prepared to patch that resource. The user must manually edit/change the resource (i.e. `oc edit <kind>/<name>`; search for the source storage class name; replace with appropriate target storage class name).

Example of "manual" action:

```
-----
Time: 2023-05-31T06:05:39.575+0000 level=info - patch-resource:
cpdinstance, configmaps, db2oltp-1685056932270995-db2oltp-cm,
[{"action":"manual","type":"string","sourcepath":".data.genkeys.sh","sourc
evalue":"--source-file-storage-class","targetvalue":"--target-file-
storage-class"}]
```

```
Time: 2023-05-31T06:05:41.176+0000 level=info - Migrate Action:
{"action":"manual","type":"string","sourcepath":".data.genkeys.sh","source
value":"--source-file-storage-class","targetvalue":"--target-file-storage-
class"}
```

```
Time: 2023-05-31T06:05:41.492+0000 level=warning - Manual Edit Required
.data.genkeys.sh: "cpd-storage"
```

```
Time: 2023-05-31T06:05:41.494+0000 level=info - configmaps db2oltp-
1685056932270995-db2oltp-cm - No Change
-----
```

Command (assumes new storage classes are default ODF storage class names):

```
$ ./cpd-migrate.sh restore \
--input-file cpd-migrate.input.json \
--cpd-operand-namespace <cpd-namespace> \
--source-block-storage-class <source-cluster-nfs-sc> \
--target-block-storage-class ocs-storagecluster-ceph-rbd \
--source-file-storage-class <source-cluster-nfs-sc> \
--target-file-storage-class ocs-storagecluster-cephfs \
--output-directory /tmp/cpd-migrate
```

### Validate patching completed successfully

After the patching activities, validate that the cpd namespace has no more references to the source cluster's storage class. With the exception of the resources that the cpd-migrate.sh tool identified as "ignore" or "manual", which need to be handled manually.

Run the following command on the target cluster to query for any remaining references to the old storage class, and handle appropriately:

```
NSPACE="<cpd-namespace>";SOURCE_SC="<source-storage-class>";for cr in $(oc
get $(oc api-resources --namespaced=true --verbs=list -o name | grep -vE
"packagemanifests|events|endpointslice|persistentvolumeclaim" | awk
'{printf "%s%s",sep,$0;sep=","}') --ignore-not-found -n ${NSPACE} -o
name);do cr_grep=$(oc get $cr -oyaml |grep -E "${SOURCE_SC}" |grep -v "
{"apiVersion":"\":"");if [[ -n "$cr_grep" ]];then echo $cr;fi;done
```

## Start Cloud Pak for Data instance on target cluster

### Prepare target cluster for Cloud Pak for Data

Before resuming CPD and services on the target cluster, make sure it has been prepared for CPD. Create custom SCCs and change required node settings. Reference: <https://www.ibm.com/docs/en/cloud-paks/cp-data/4.6.x?topic=installing-preparing-your-cluster>

Examples:

- Db2 kubeletconfig: `cpd-cli manage apply-db2-kubelet --openshift-type=self-managed`
- WKC SCC: `cpd-cli manage apply-scc --cpd_instance_ns=<cpd-namespace> --components=wkc`
- CRI-O settings: `cpd-cli manage apply-crio --openshift-type=self-managed`

Start services on the target cluster that were manually shutdown on source cluster

For the services and resources that were manually shutdown in the **"Scale down remaining services and resources running in CPD namespace"** step, startup those services now.

- Examples: nginx deployment, db2ucluster statefulsets, foundationdb, MANTA, openscale, etc.

#### nginx deployment (scale back up to 2):

```
oc scale --replicas=2 deploy/ibm-nginx -n <cpd-namespace>
```

#### MANTA scale up replicas (back to 1):

```
oc patch -p '{"spec":{"replicas":1}}' --type=merge mantaflow mantaflow-wkc
-n <cpd-namespace>
```

#### Openscale startup:

```
oc patch -p '{"spec":{"shutdown":"false"}}' --type=merge woservice
aiopenscale -n <cpd-namespace>
oc patch -p '{"spec":{"ignoreForMaintenance": false}}' --type=merge
woservice aiopenscale -n <cpd-namespace>
```

#### Foundationdb startup:

```
oc patch -p '{"spec":{"shutdown":"false"}}' --type=merge fdbcluster wkc-
foundationdb-cluster -n <cpd-namespace>
```

#### db2ucluster statefulsets (scale back up to 1):

```
oc scale --replicas=1 sts/c-db2oltp-iis-db2u -n <cpd-namespace>
oc scale --replicas=1 sts/c-db2oltp-wkc-db2u -n <cpd-namespace>

oc scale --replicas=1 sts/<db2wh-db2u-sts> -n <cpd-namespace>
oc scale --replicas=1 sts/<db2wh-etc-sts> -n <cpd-namespace>
```

Note: If there were additional db2u related statefulsets that were scaled down, scale them up here to the same replica number as it was on the source cluster.

#### Datastage deploy and statefulset

```
## Record this output for the resume operations:
oc get deploy,sts -n <cpd-namespace>|grep ds-px-default
```

```
oc scale --replicas=1 deploy/ds-px-default-ibm-datastage-px-runtime -n
<cpd-namespace>
oc scale --replicas=1 sts/ds-px-default-ibm-datastage-px-compute -n <cpd-
namespace>
```

**Remaining resources:** If there were additional services, deployments, statefulsets, or other that were manually shutdown on the source cluster, resume them here.

### Examples of starting the resources:

#### Scale up nginx deployment

```
$ oc get deploy -n cpd |grep -E "NAME|nginx"
NAME                                     READY    UP-TO-DATE
AVAILABLE    AGE
aiopenscale-ibm-aio-nginx               0/0      0          0
2d13h
ibm-nginx                               0/0      0          0
2d13h
ibm-nginx-tester                        0/0      0          0
2d13h
spark-hb-nginx                          0/0      0          0
2d13h
$ oc scale --replicas=2 deploy/ibm-nginx -n cpd
deployment.apps/ibm-nginx scaled
$ oc get deploy ibm-nginx -n cpd
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
ibm-nginx     2/2      2             2            2d13h
$ oc get pod -n cpd
NAME                                     READY    STATUS    RESTARTS    AGE
ibm-nginx-7cd879b767-fjgn4             1/1      Running   0           4m47s
ibm-nginx-7cd879b767-fn22k             1/1      Running   0           4m47s
```

#### Scale up db2u statefulsets

```
$ oc -n cpd get sts -n cpd |grep -E "NAME|db2"
NAME                                     READY    AGE
c-db2oltp-iis-db2u                     0/0      2d18h
c-db2oltp-wkc-db2u                     0/0      8d
c-db2wh-1686277435040922-db2u         0/0      15h
c-db2wh-1686277435040922-etcd         0/0      15h

$ oc scale --replicas=1 sts/c-db2oltp-wkc-db2u -n cpd
statefulset.apps/c-db2oltp-wkc-db2u scaled
$ oc scale --replicas=1 sts/c-db2oltp-iis-db2u -n cpd
statefulset.apps/c-db2oltp-iis-db2u scaled
$ oc scale --replicas=1 sts/c-db2wh-1686277435040922-db2u -n cpd
statefulset.apps/c-db2wh-1686277435040922-db2u scaled
```

```
$ oc scale --replicas=1 sts/c-db2wh-1686277435040922-etcd -n cpd
statefulset.apps/c-db2wh-1686277435040922-etcd scaled

$ oc -n cpd get pods|grep db2
c-db2oltp-iis-db2u-0                                1/1      Running
0                                                    44h
c-db2oltp-iis-instdb-sx2hm                          0/1
Completed 0                                         2d19h
c-db2oltp-wkc-db2u-0                                1/1      Running
0                                                    44h
c-db2wh-1686277435040922-db2u-0                    1/1      Running
0                                                    15h
c-db2wh-1686277435040922-etcd-0                    1/1      Running
0                                                    15h
```

## Start foundationdb

Check foundationdb CR state from CPD namespace:

```
$ oc -n cpd get fdbcluster wkc-foundationdb-cluster -o yaml |grep -E
"ceph|shutdown"

    pvcStorageClass: ocs-storagecluster-cephfs
      storageClassName: ocs-storagecluster-ceph-rbd
      storageClassName: ocs-storagecluster-ceph-rbd
shutdown: "true"
shutdownStatus: shutdown
```

Start foundationdb-cluster CR from CPD namespace:

```
$ oc -n cpd get fdbcluster wkc-foundationdb-cluster -o json |jq
'.spec.shutdown'
"true"
$ oc -n cpd edit fdbcluster wkc-foundationdb-cluster
fdbcluster.foundationdb.opencontent.ibm.com/wkc-foundationdb-cluster
edited
$ oc -n cpd get fdbcluster wkc-foundationdb-cluster -o json |jq
'.spec.shutdown'
"false"
$

$ oc get pods -n cpd
NAME                                READY   STATUS    RESTARTS
AGE
c-db2oltp-wkc-db2u-0                1/1     Running   0
35m
ibm-nginx-7cd879b767-fjgn4          1/1     Running   0
45m
ibm-nginx-7cd879b767-fn22k          1/1     Running   0
```

```

45m
wkc-foundationdb-cluster-cluster-controller-1  2/2    Running    0
18m
wkc-foundationdb-cluster-log-1                2/2    Running    0
18m
wkc-foundationdb-cluster-log-2                2/2    Running    0
18m
wkc-foundationdb-cluster-log-3                2/2    Running    0
18m
wkc-foundationdb-cluster-log-4                2/2    Running    0
18m
wkc-foundationdb-cluster-storage-1            2/2    Running    0
18m
wkc-foundationdb-cluster-storage-2            2/2    Running    0
18m
wkc-foundationdb-cluster-storage-3            2/2    Running    0
18m
wkc-foundationdb-cluster-storage-4            2/2    Running    0
18m
wkc-foundationdb-cluster-storage-5            2/2    Running    0
18m
wkc-foundationdb-cluster-storage-6            2/2    Running    0
18m

```

## Start MANTA

From CPD namespace, set replicas to 1:

```

$ oc -n cpd get mantaflow mantaflow-wkc -o json | jq '.spec.replicas'
0
$ oc -n cpd edit mantaflow mantaflow-wkc
mantaflow.adl.getmanta.com/mantaflow-wkc edited
$ oc -n cpd get mantaflow mantaflow-wkc -o json | jq '.spec.replicas'
1

```

## Start Openscale:

From CPD namespace, set shutdown=false in the two openscale CR instances:

```

$ oc patch W0Service openscale-defaultinstance -n cpd --type merge --
patch '{"spec": {"shutdown":"false"}}'
woservice.wos.cpd.ibm.com/openscale-defaultinstance patched
$ oc -n cpd patch W0Service aiopenscale -n cpd --type merge --patch
'{"spec": {"shutdown": "false"}}'
woservice.wos.cpd.ibm.com/aiopenscale patched
$ oc -n cpd get W0Service openscale-defaultinstance -o json | jq
'.spec.shutdown'
>false"
$ oc -n cpd get W0Service aiopenscale -o json | jq '.spec.shutdown'

```

```
"false"  
$
```

## Datastage

```
$ oc scale --replicas=1 deploy/ds-px-default-ibm-datastage-px-runtime -n  
cpd  
deploy.apps/ds-px-default-ibm-datastage-px-runtime scaled  
$ oc scale --replicas=1 sts/ds-px-default-ibm-datastage-px-compute -n cpd  
statefulset.apps/ds-px-default-ibm-datastage-px-compute scaled
```

## Resume CPD service operations with restore posthooks

The cpd-cli oadp restore posthooks will take all of the CPD services out of maintenance mode and scale back up all of the deployments, statefulsets, etc.

Example Command:

```
cpd-cli oadp restore posthooks --include-namespaces <cpd-namespace> --log-  
level=debug --verbose --scale-wait-timeout
```

Run the following commands to verify CPD instance is running:

```
$ oc get pods -n <cpd-namespace>  
$ oc get route -n <cpd-namespace>  
$ oc get pvc -n <cpd-namespace>  
$ oc get sts -n <cpd-namespace>  
$ oc get deployments -n <cpd-namespace>
```

Get CR status:

```
$ cpd-cli manage get-cr-status --cpd_instance_ns=<cpd-namespace>  
$ oc get ccs -n <cpd-namespace>  
$ oc get wkc wkc-cr -o yaml -n <cpd-namespace>
```

**Validate that all installed CPD services resume to Completed status:**

```
cpd-cli manage get-cr-status --cpd_instance_ns=<cpd-namespace>
```

After the command completes successfully, the migration restore is now complete!



